Figure 0.1. Signal-to-learning operating chain

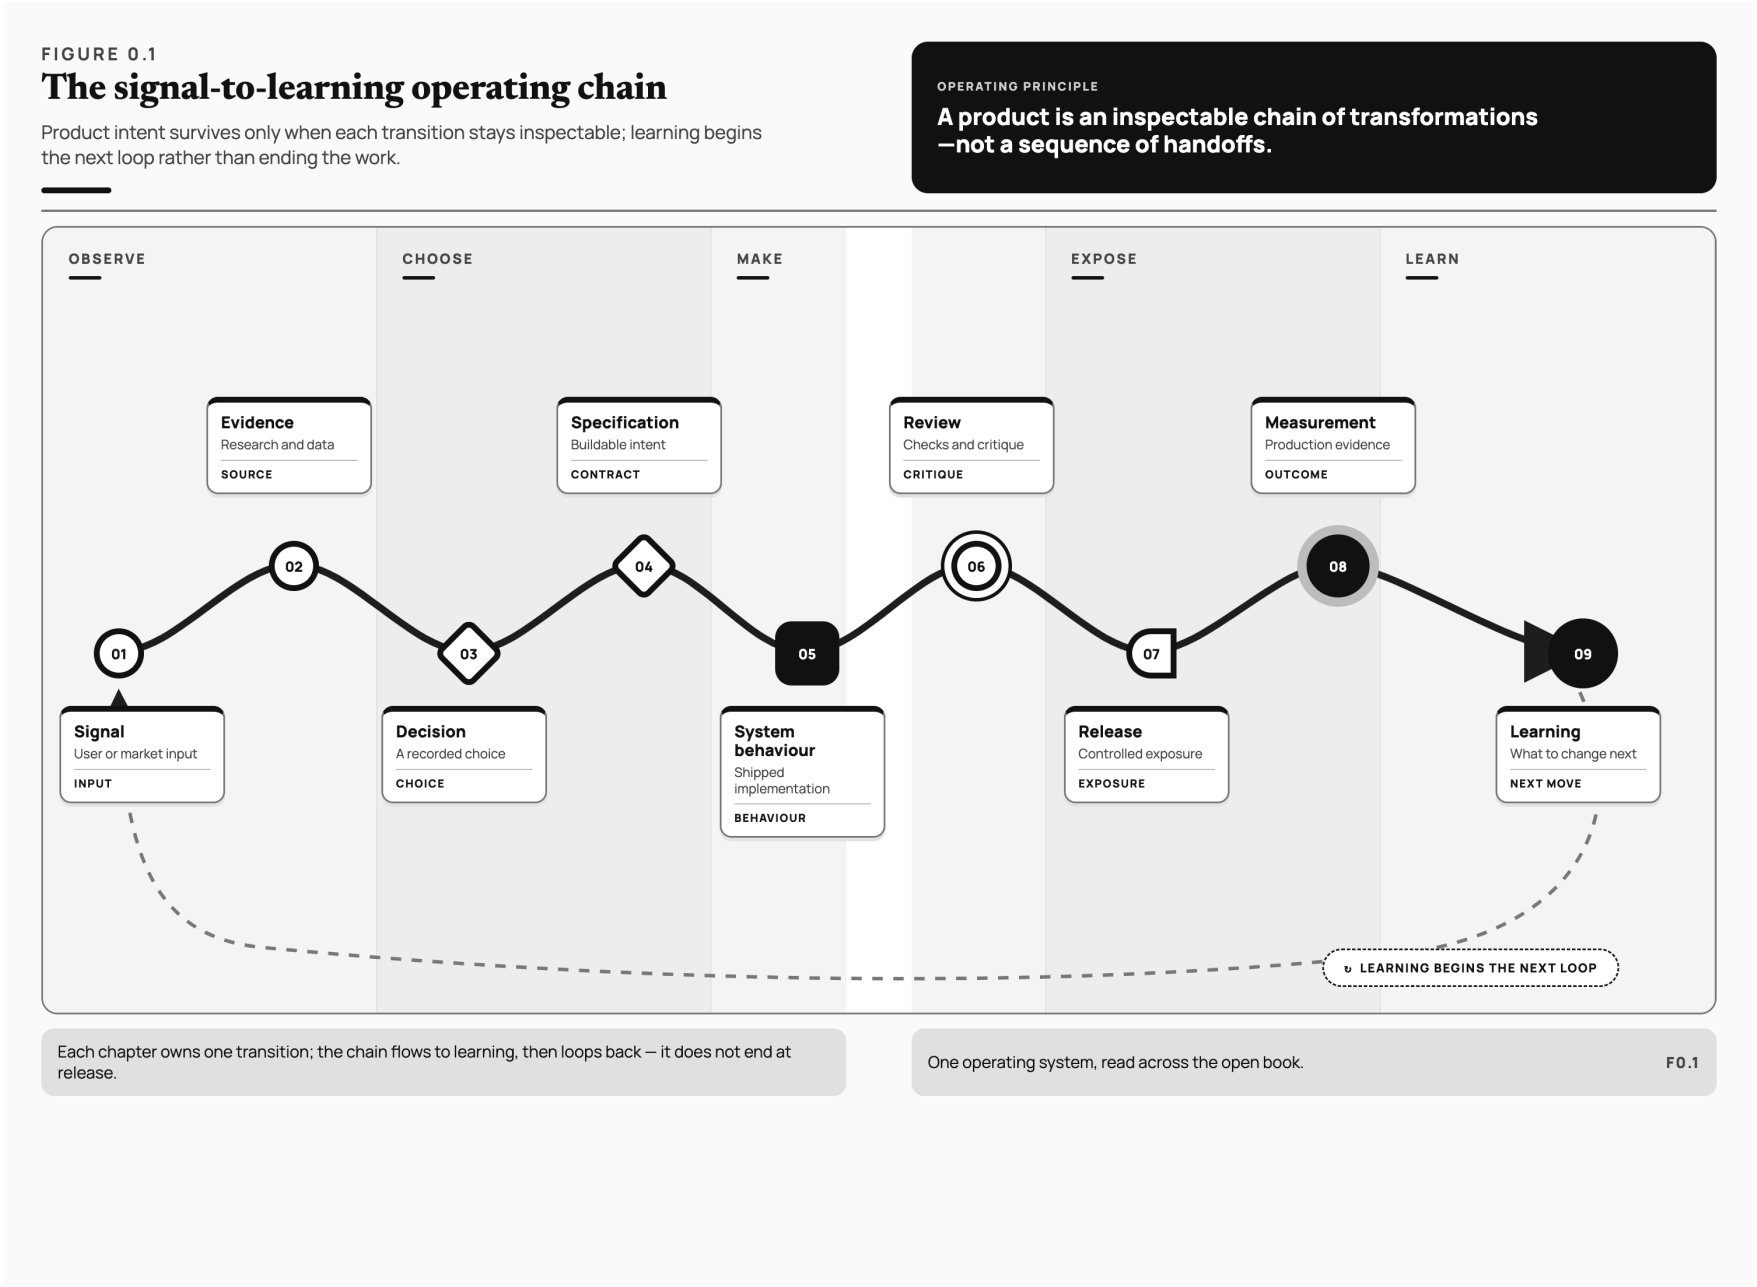

Product intent survives only when each transition remains inspectable; learning begins the next loop rather than ending the work.

Figure 1.1. Participation ladder: ask, inspect, prototype, contribute

Technical participation progresses through observable work, not status; the appropriate stopping point depends on the product and risk.

Figure 2.1. Registration request/response sequence

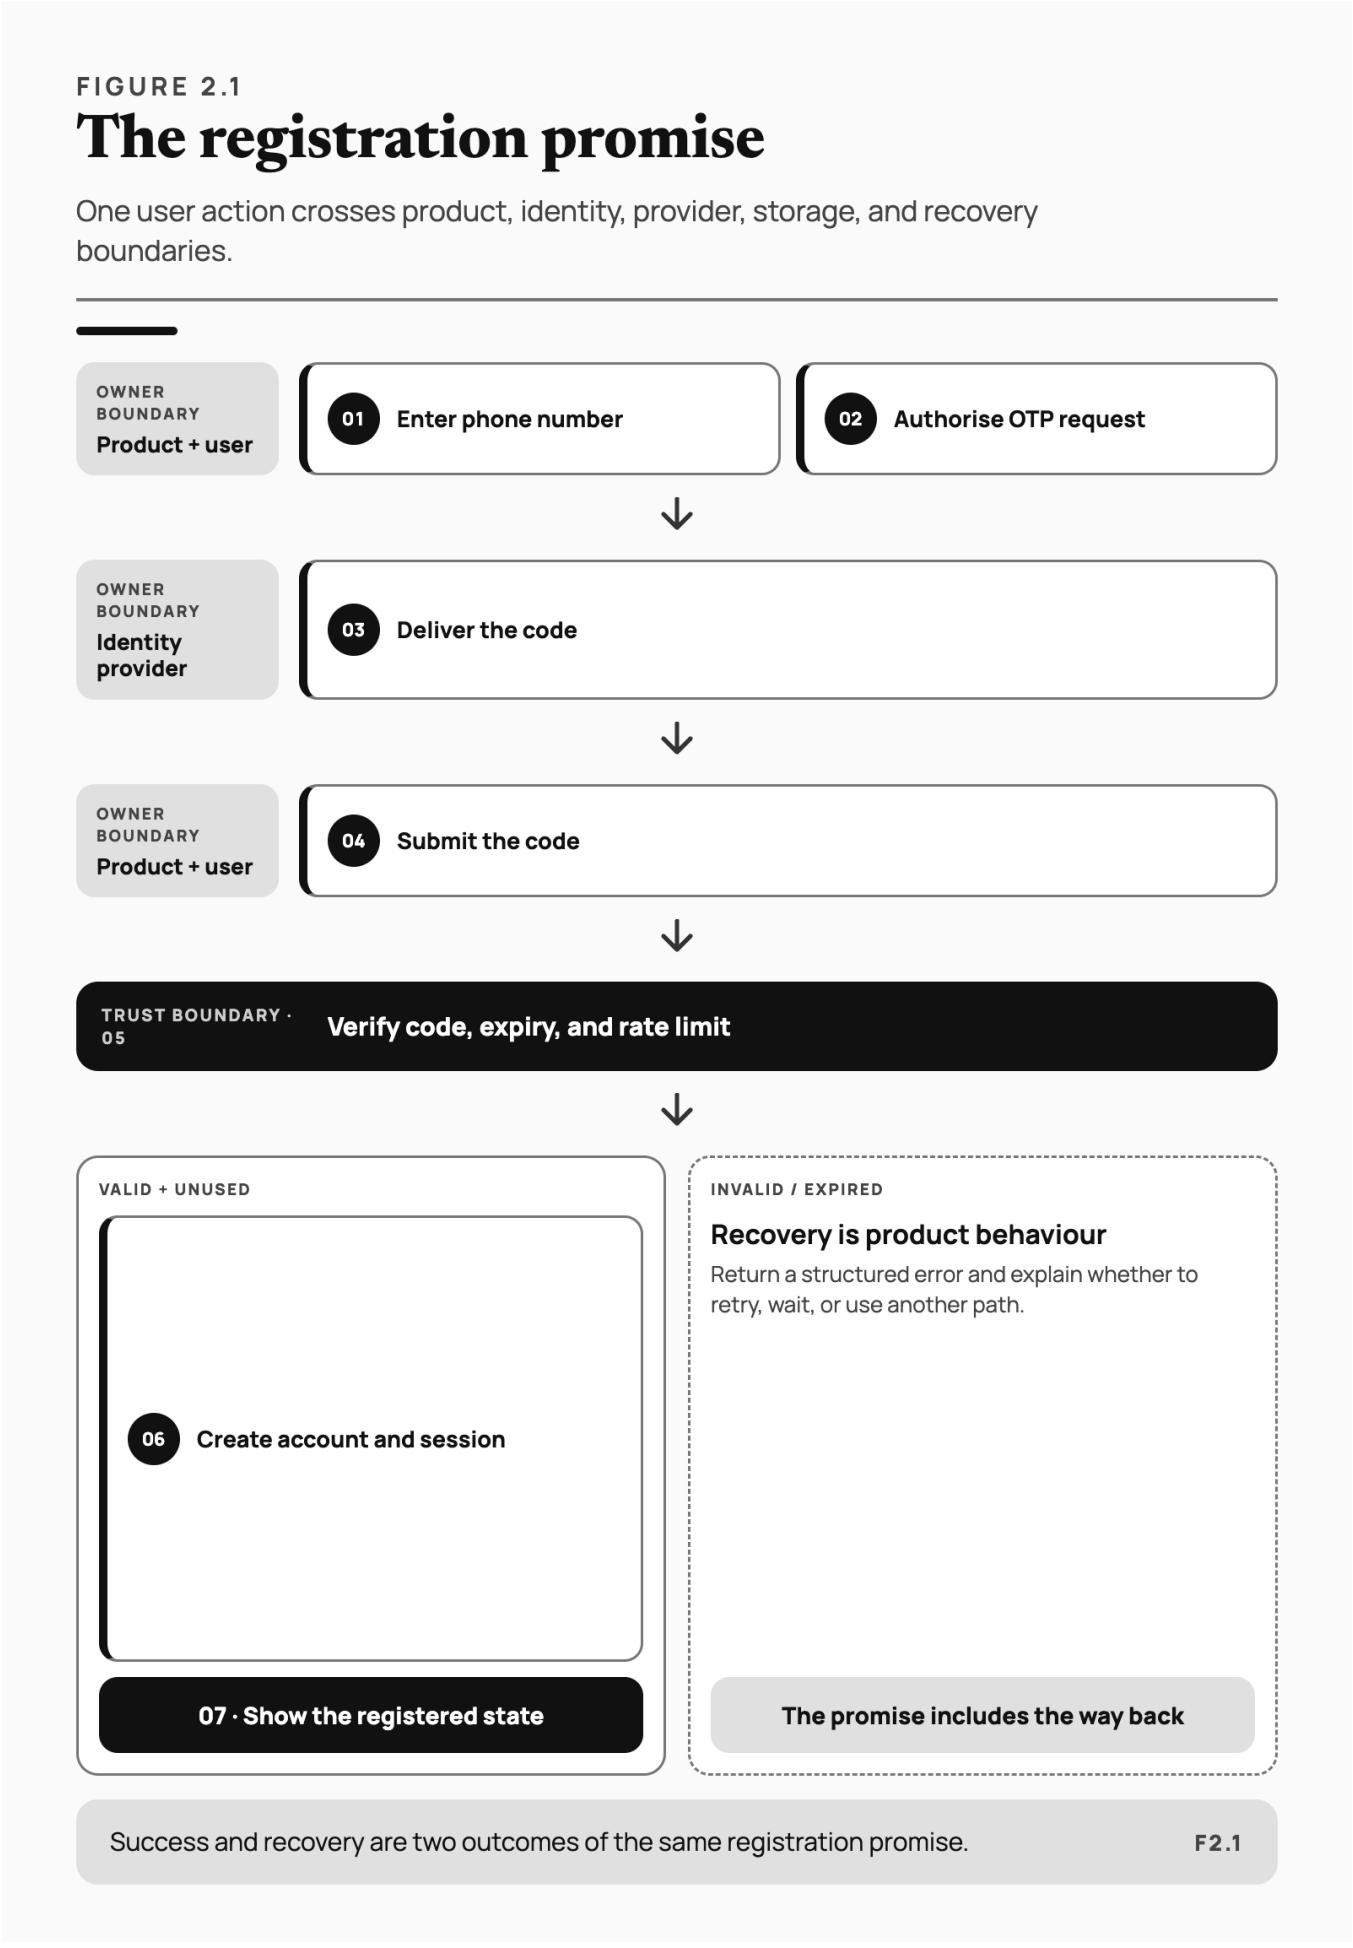

The registration promise crosses multiple actors; success and recovery states both belong in the product conversation.

Figure 3.1. Commit-to-production delivery path

FIGURE 3.1

The delivery confidence path

Integration, deployment, exposure, and learning are separate product decisions.

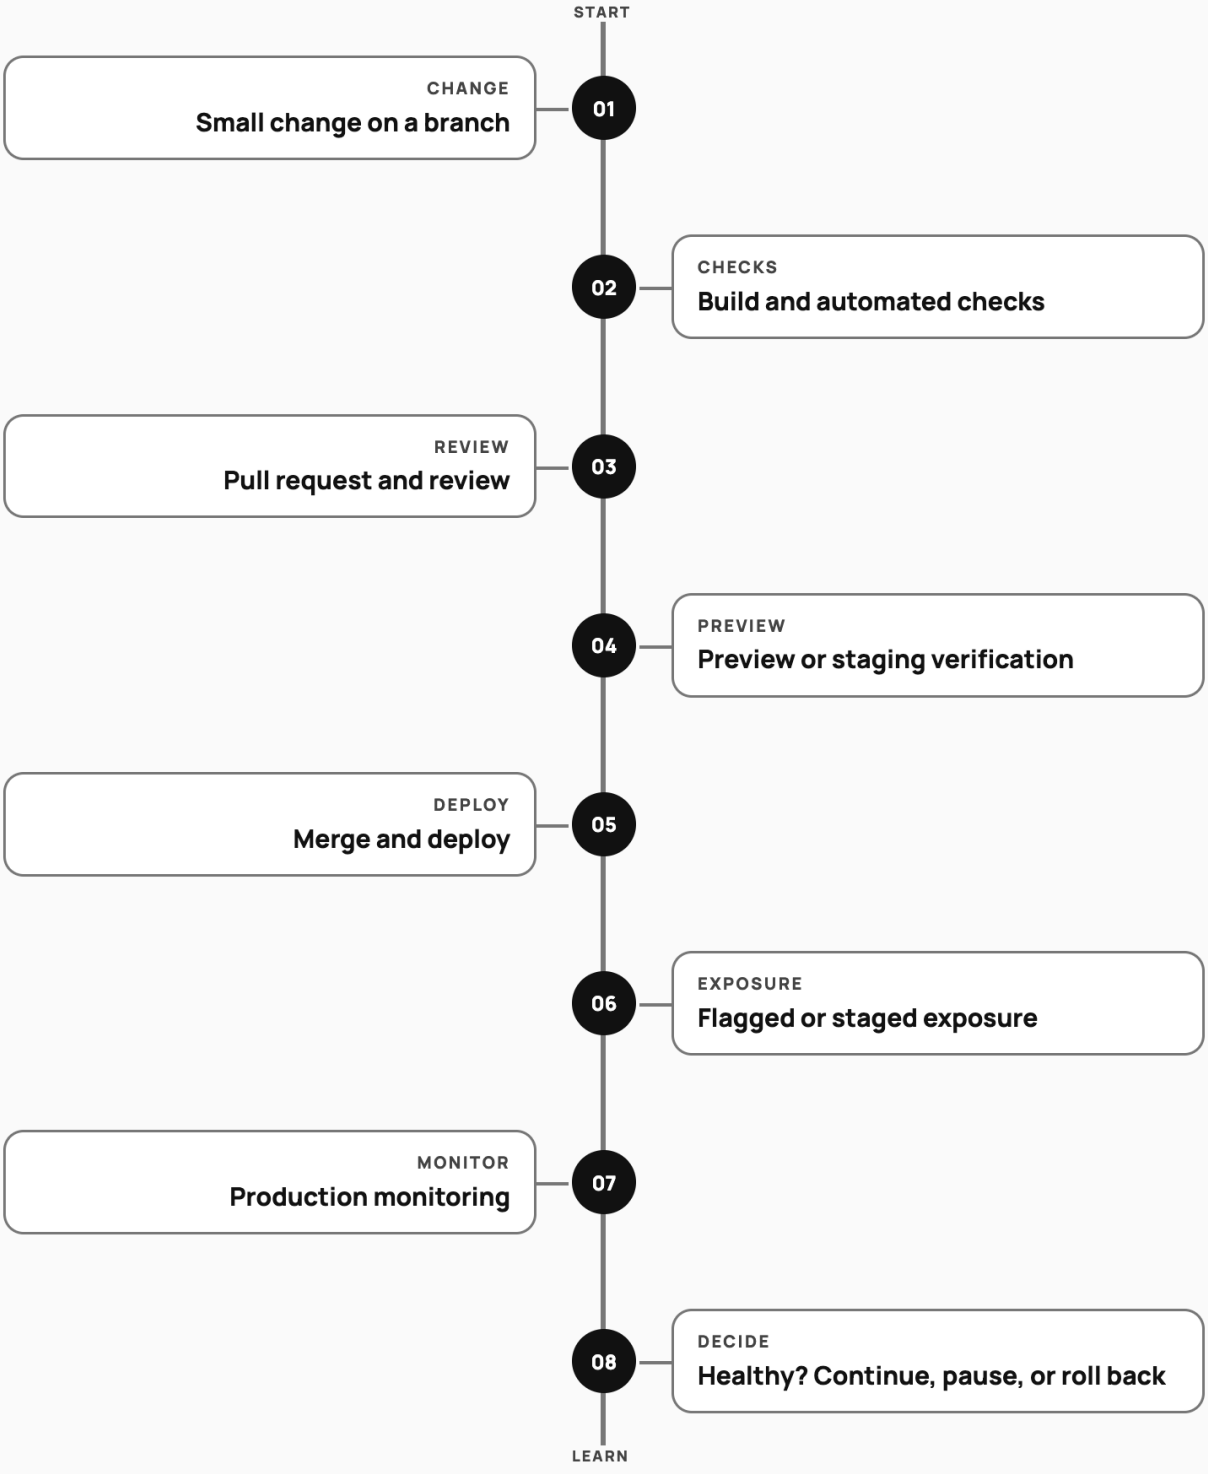

Shipping confidence accumulates one verified boundary at a time. F3.1

A delivery path separates code integration, deployment, exposure, and the decision to continue.

Figure 4.1. Original asset to responsive delivery

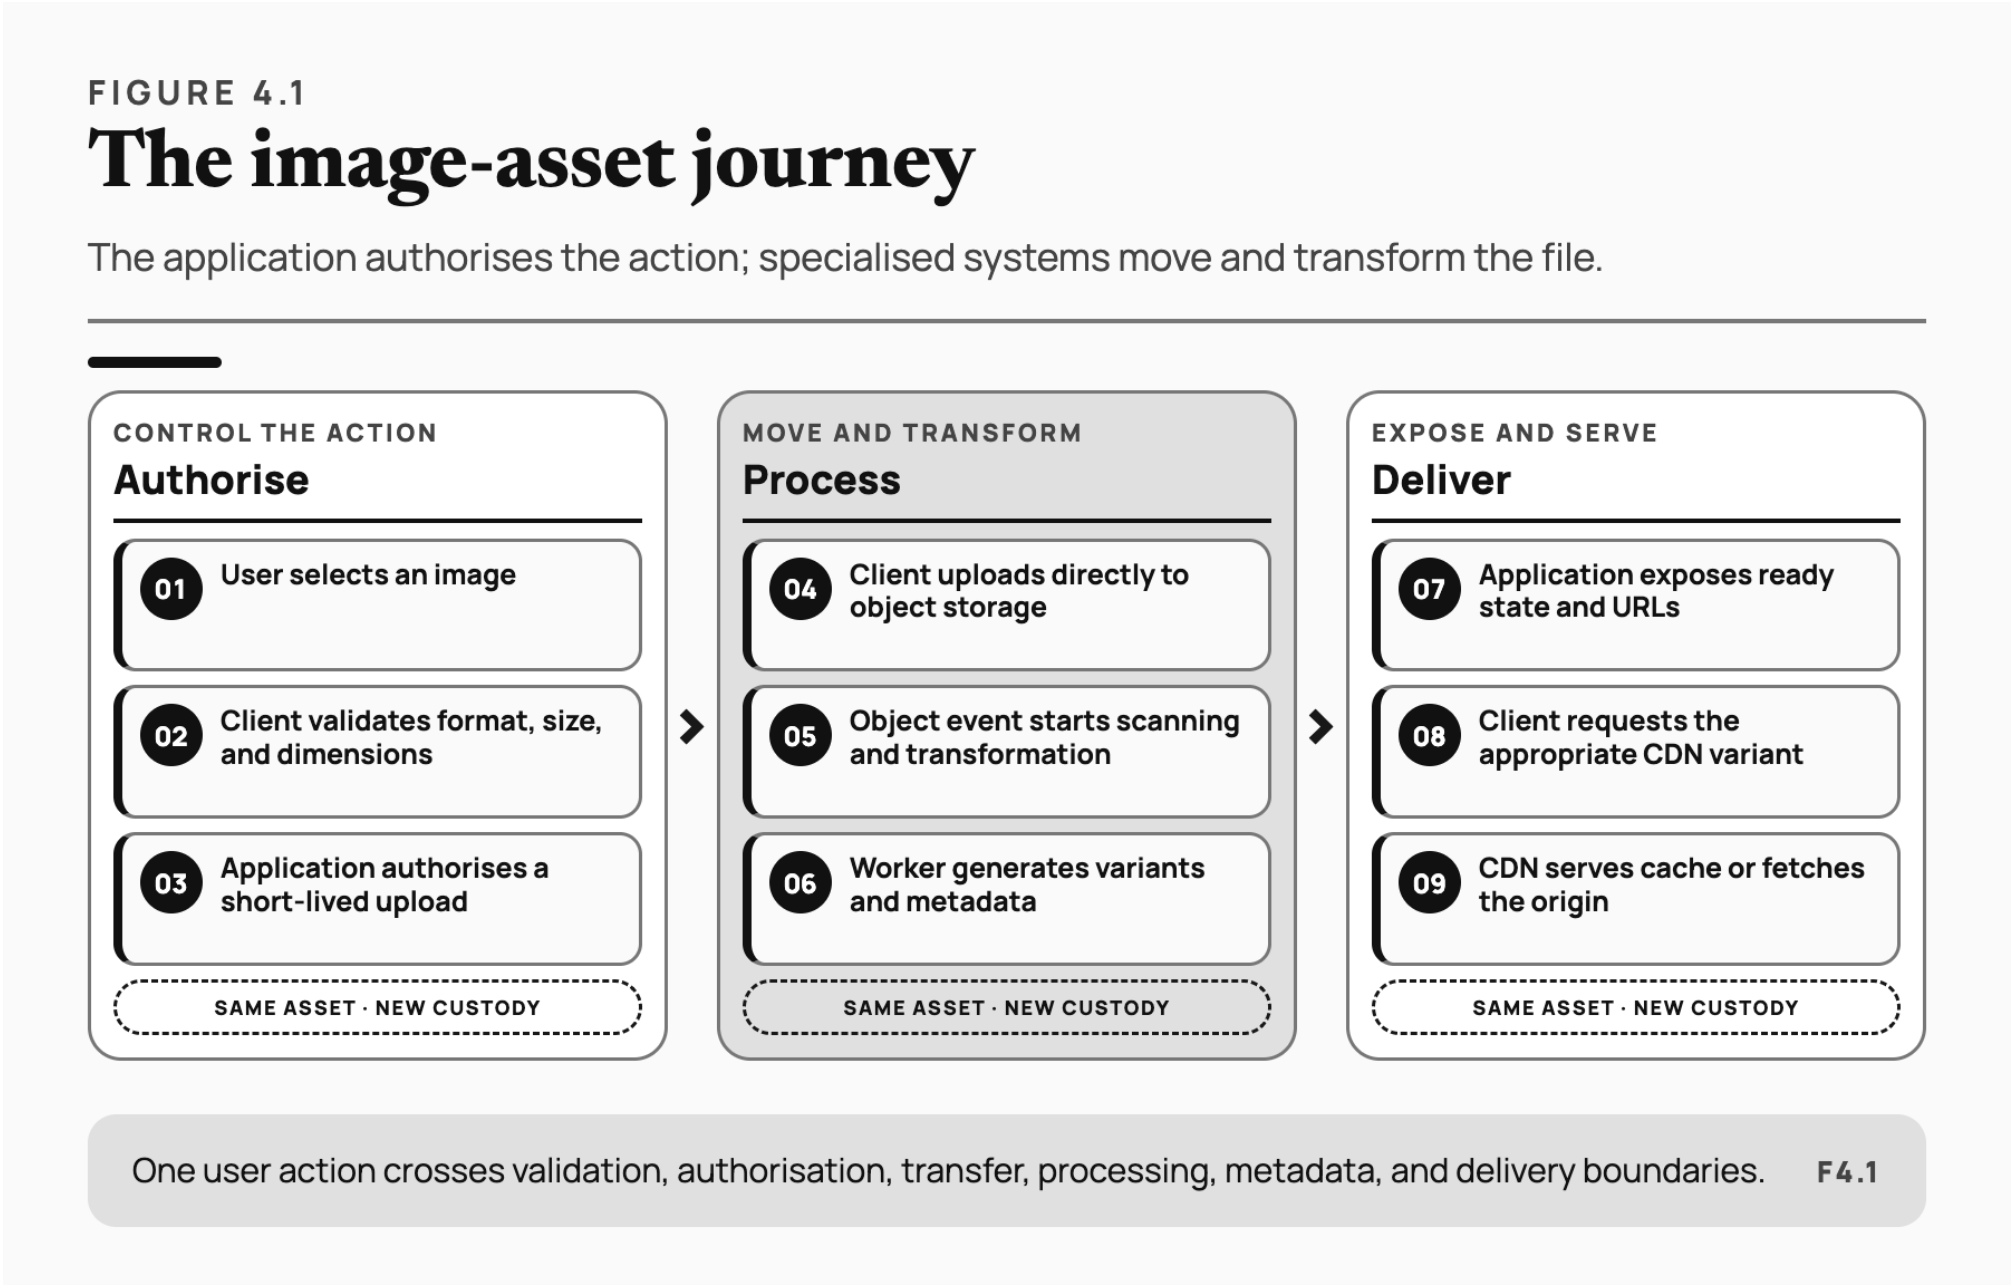

The image-asset journey connects selection and validation to authorised transfer, processing, product state, variants, and delivery. Later sections examine each boundary.

Figure 4.2. Asset-state lifecycle

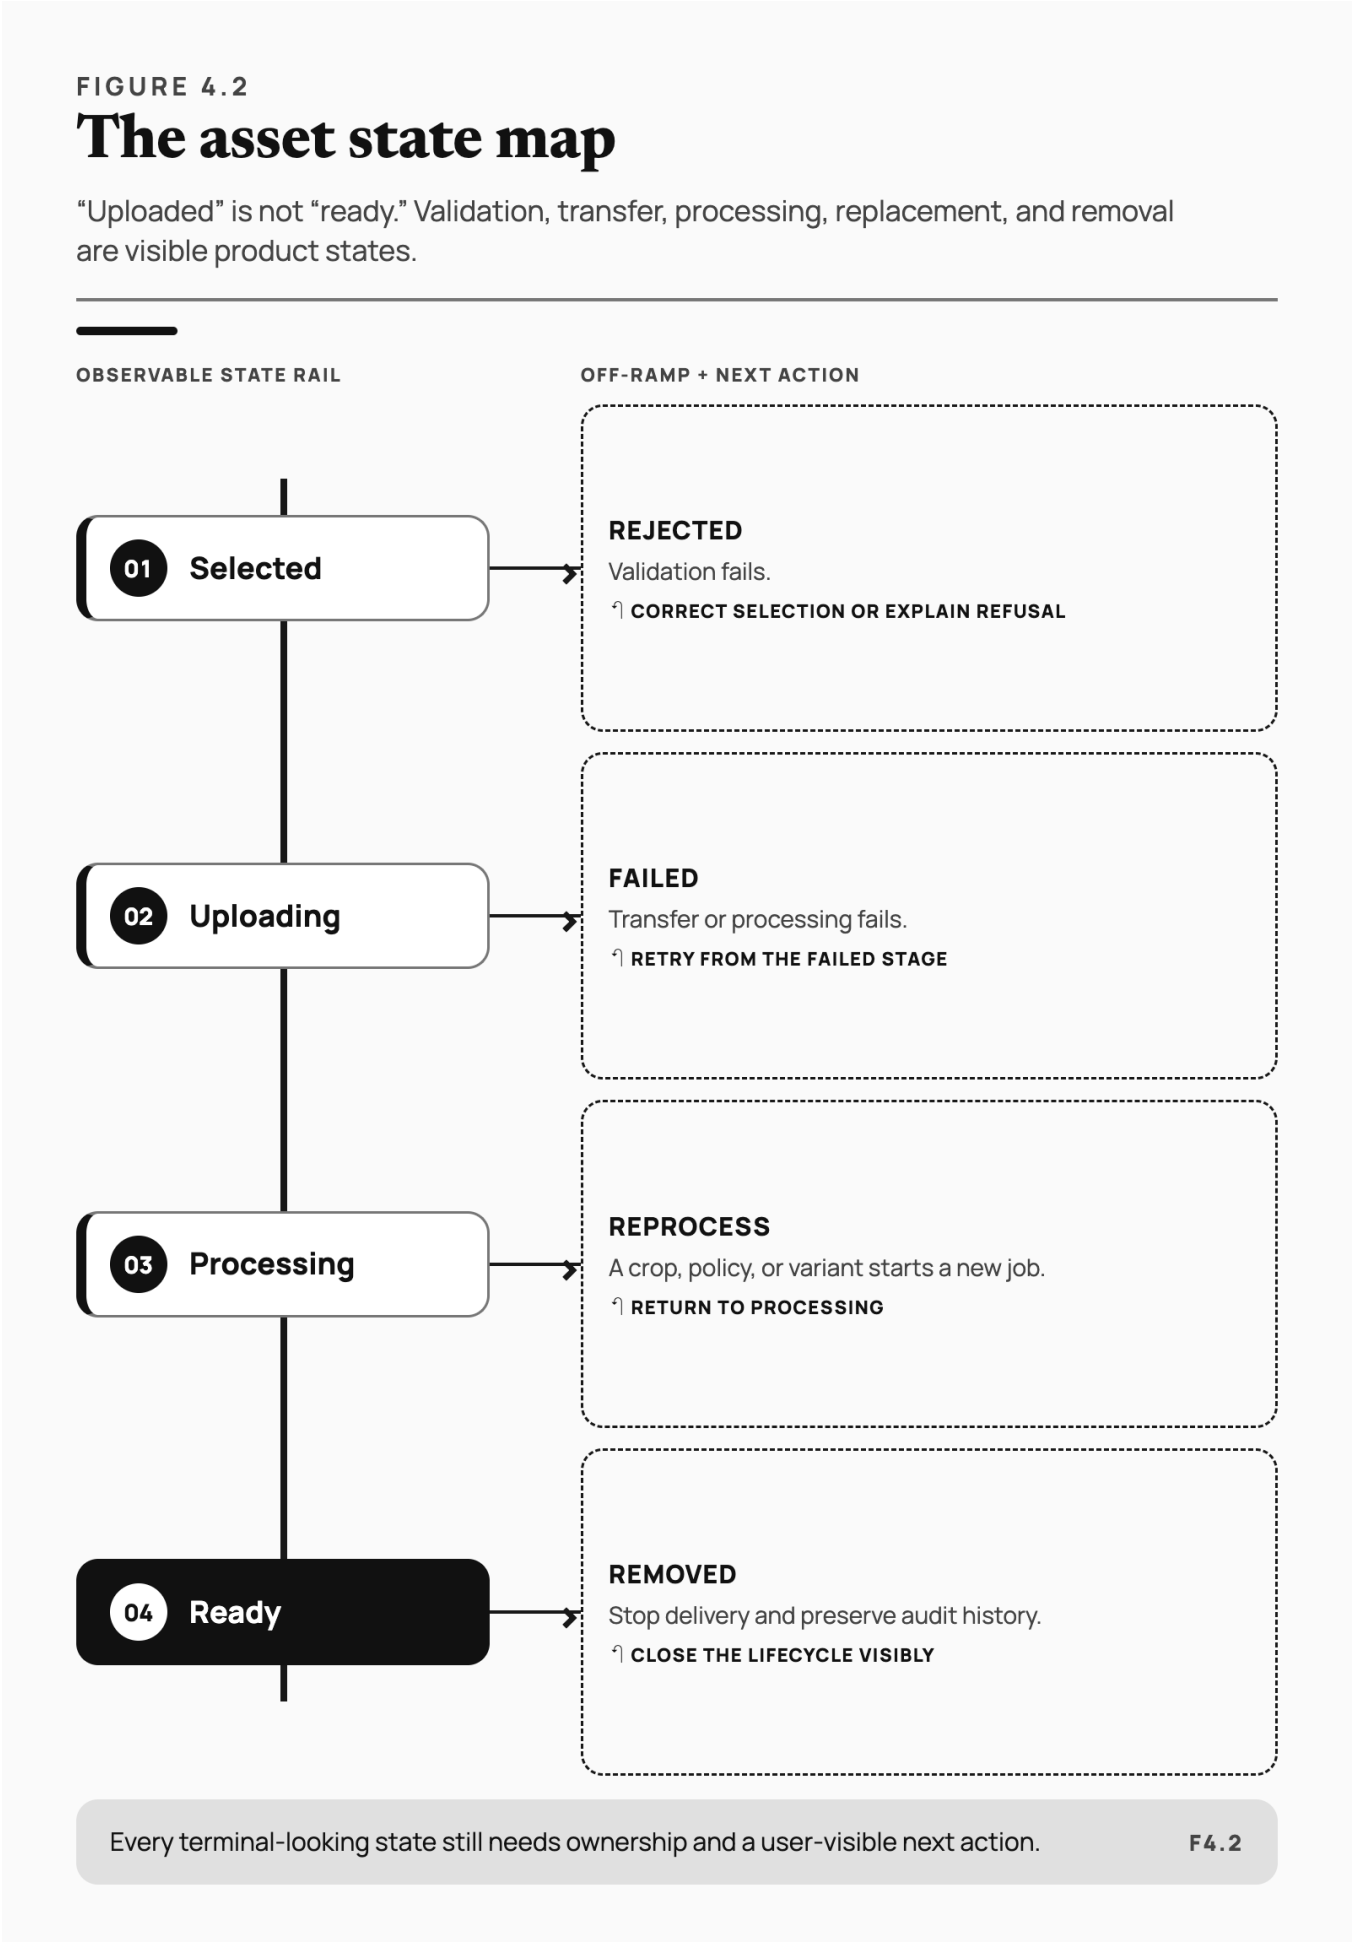

Asset state is user-visible product behavior; "uploaded" is not the only meaningful outcome.

Figure 5.1. Relational grain through joins

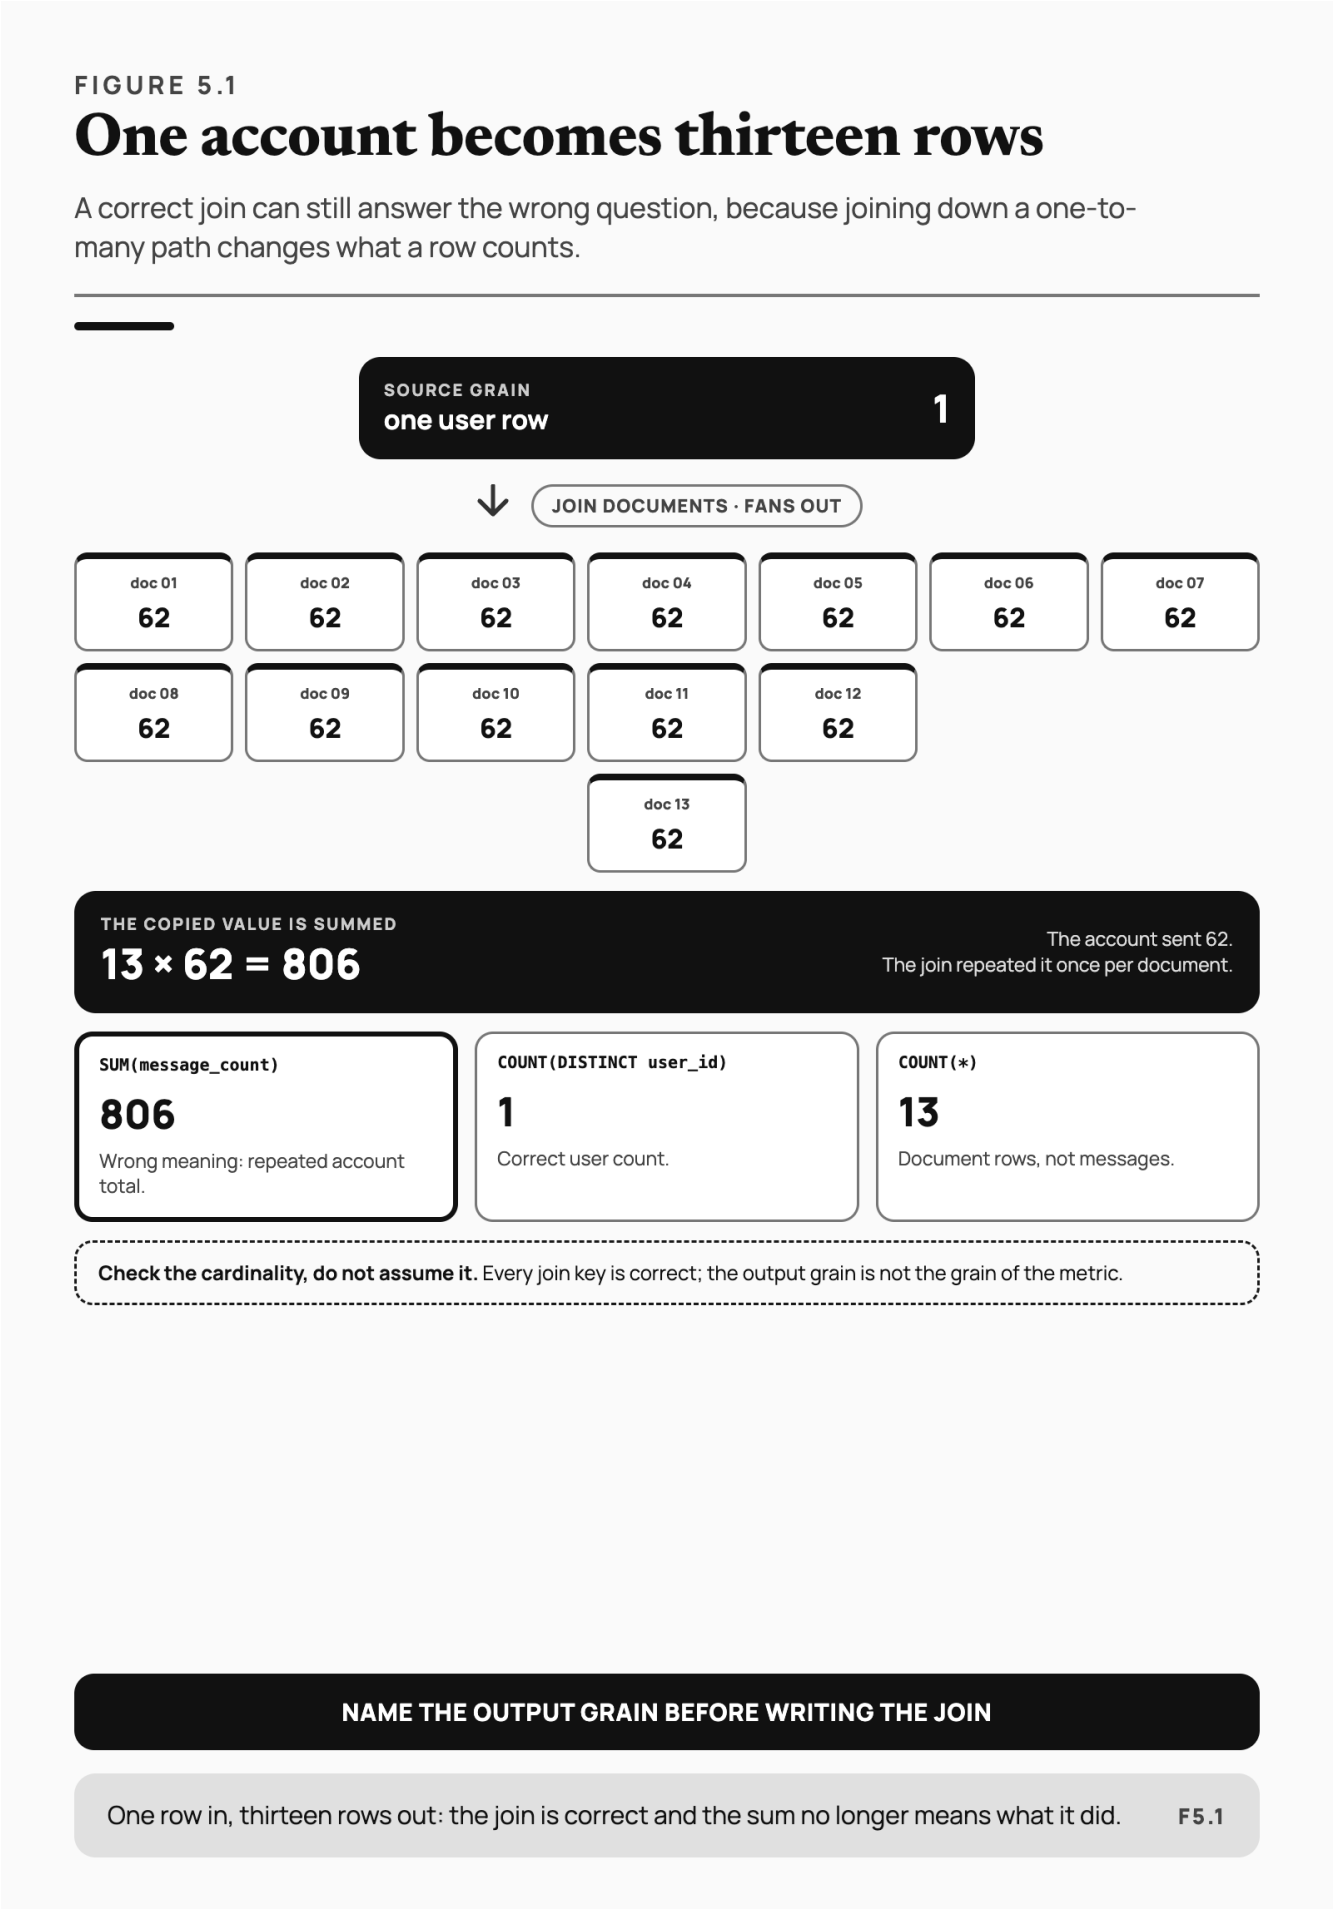

One account holding 13 documents and one message-count row returns 13 rows once both tables are joined, so summing the message column reports 806 messages instead of 62.

Figure 6.1. Event-to-metric semantic chain

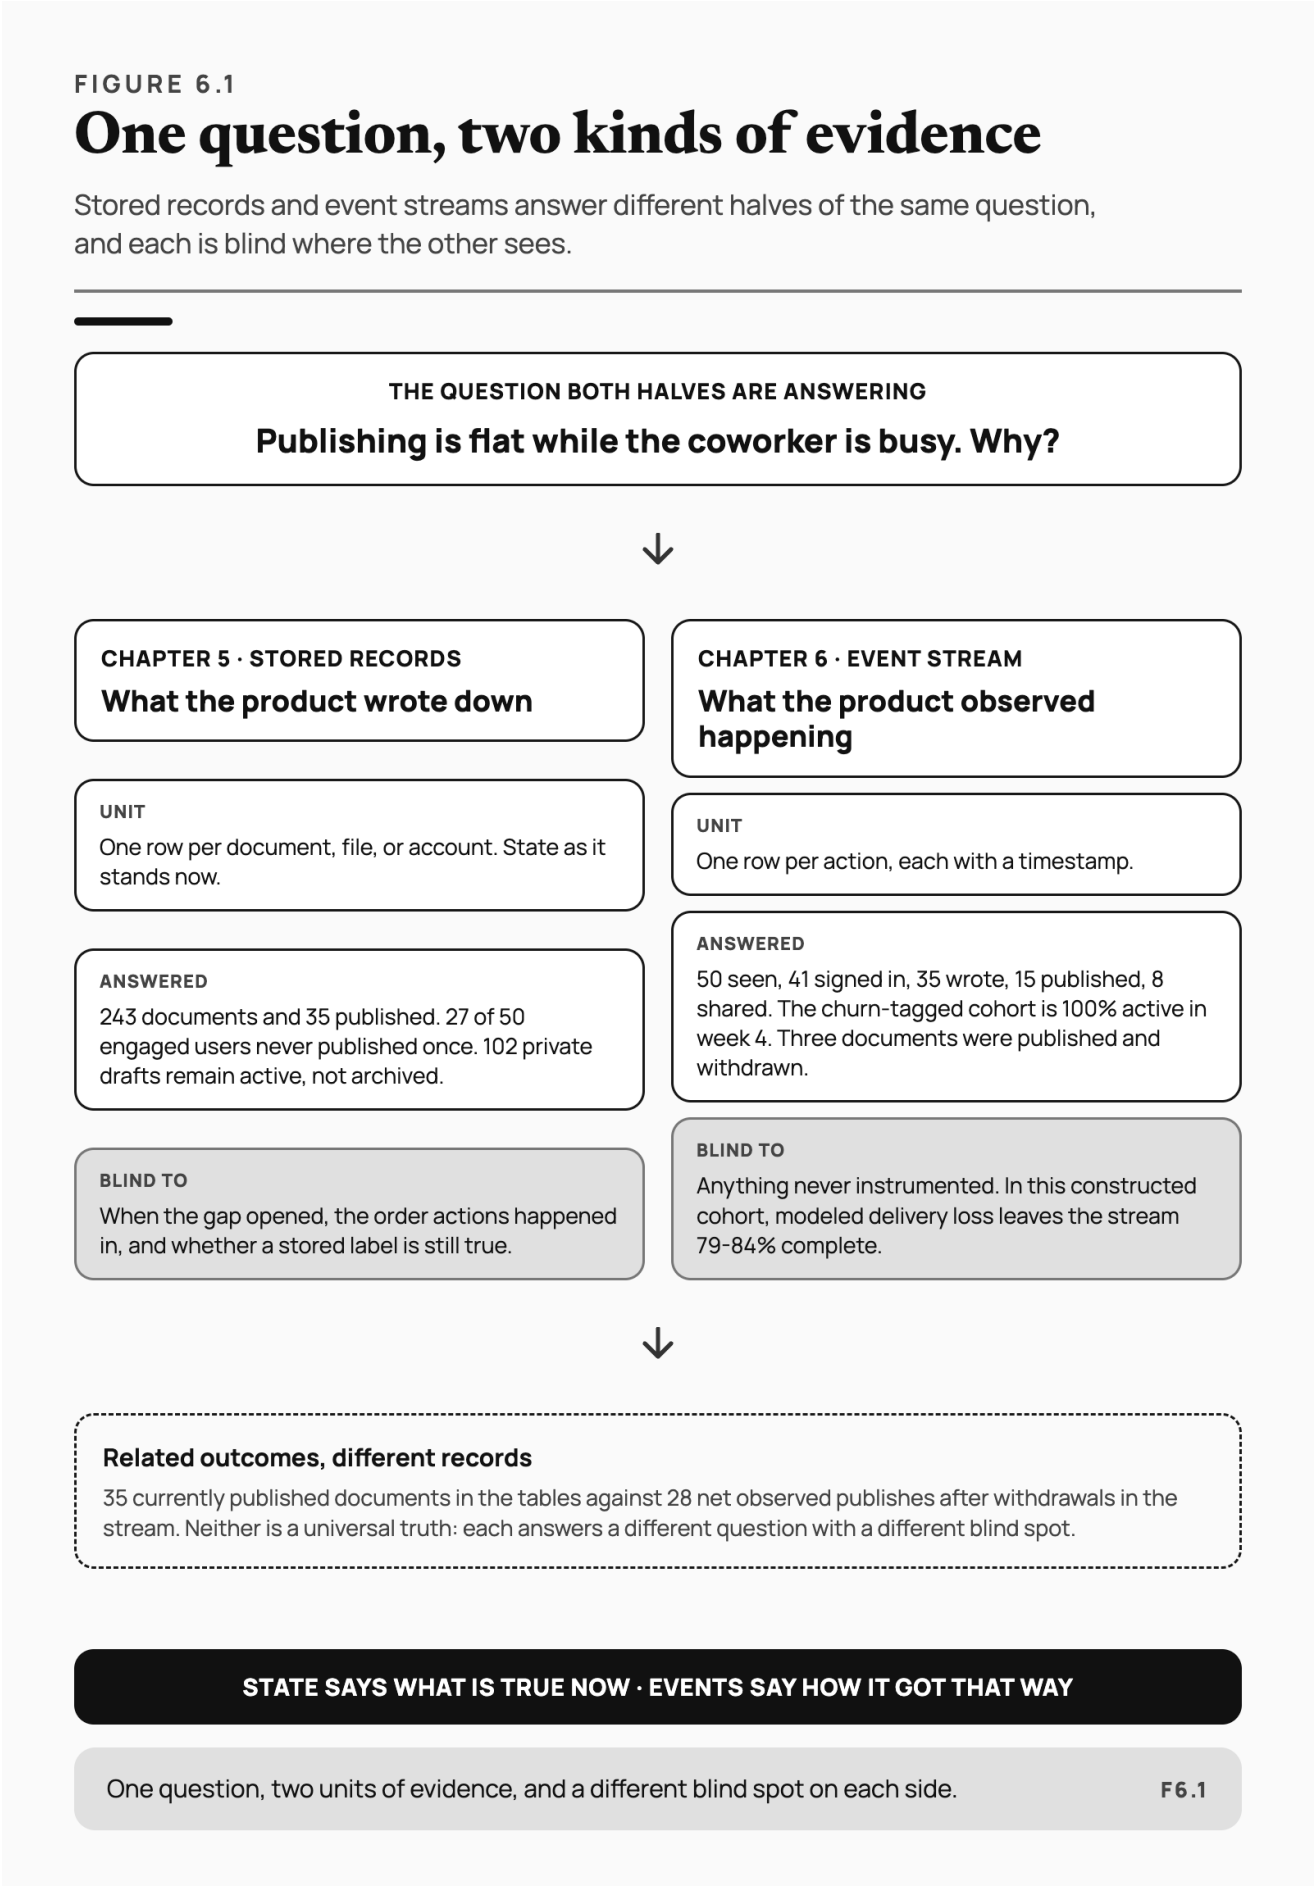

One question and two units of evidence. The tables count state as it stands and cannot say when the gap opened; the stream counts observed actions and misses whatever was never instrumented. Related publishing outcomes read 35 in current state and 28 net observed publishes after withdrawals.

Figure 7.1. Metric chain graded by evidence strength

The same four-link chain, graded link by link. Confidence falls from a randomized result to an association to an unanswered question, and the illustrative figures on the plate show the shape of that fall rather than the measured outcome.

Figure 7.2. Experiment go/no-go decision flow

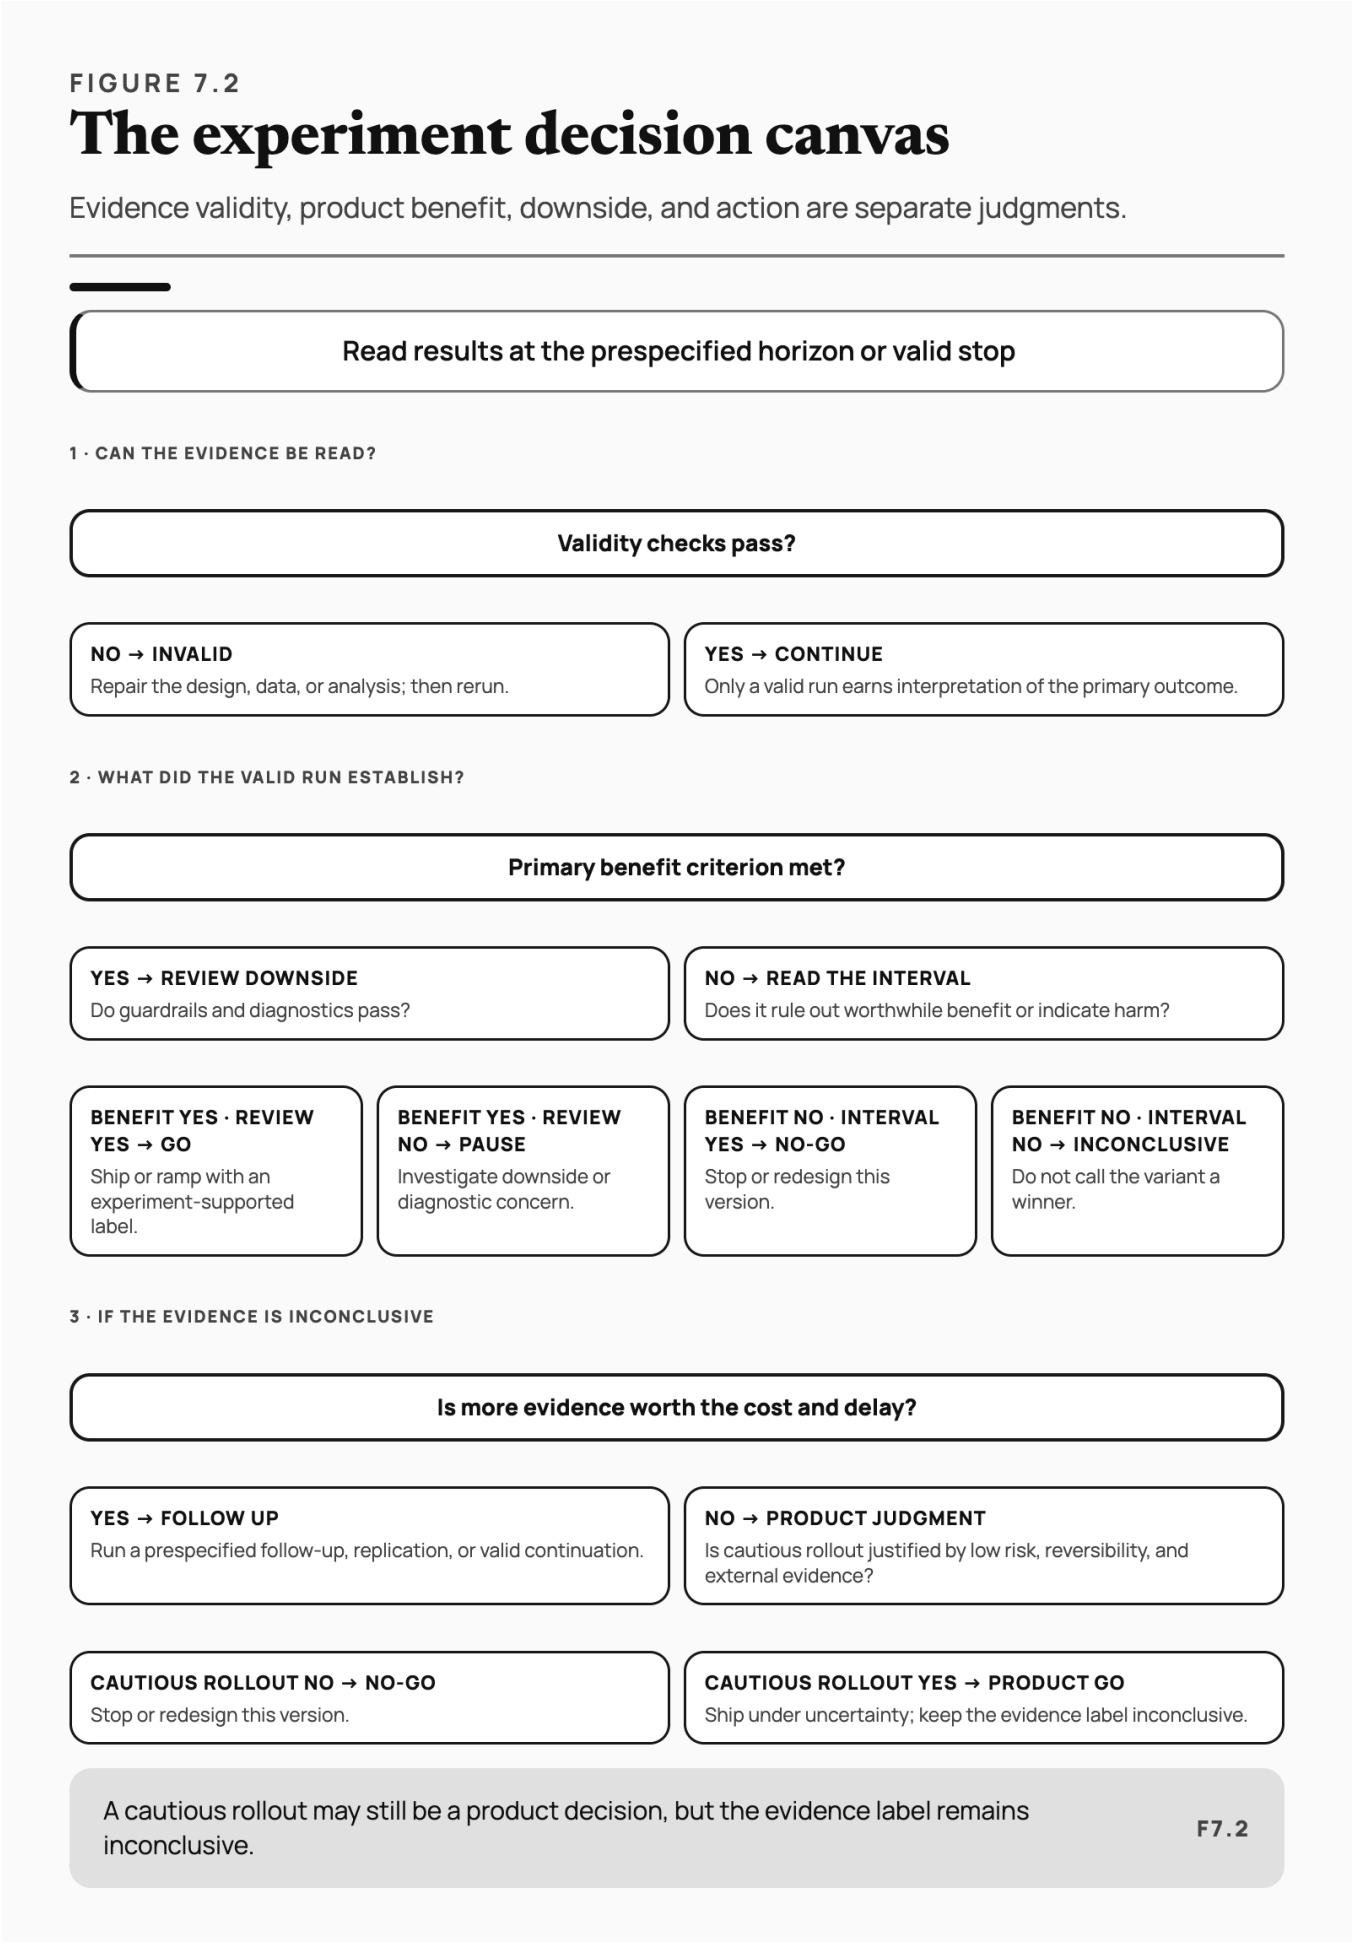

The go/no-go canvas separates evidence validity, product benefit, downside, uncertainty, and action.

Figure 8.1. Shipper reverse-logistics custody sequence

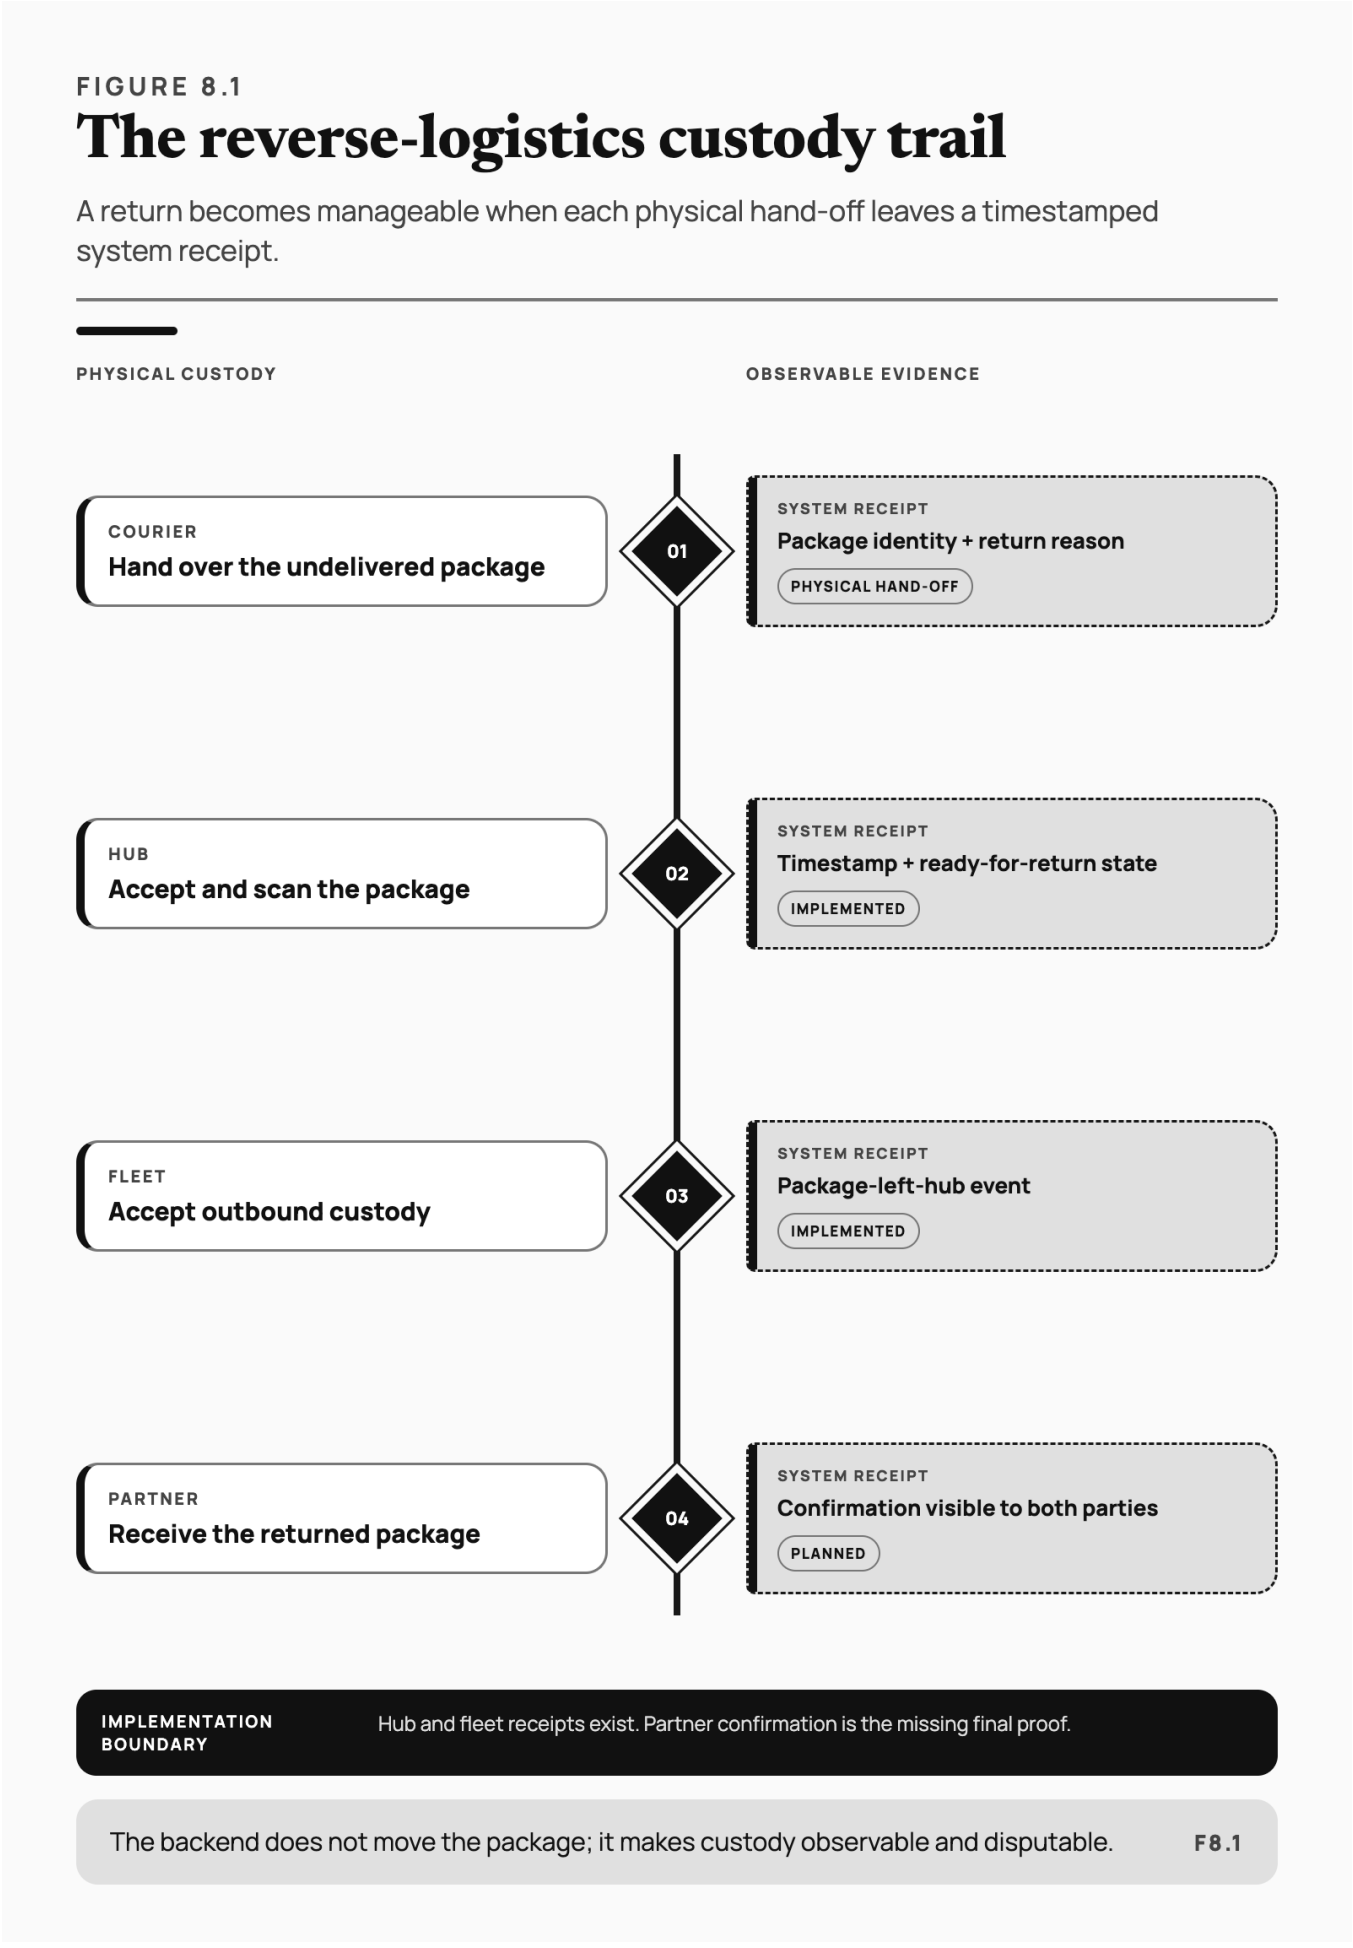

The reverse-logistics flow: hub and driver scanning were implemented, while partner-side receipt confirmation remained planned.

Figure 8.2. E-wallet flow with happy, error, and recovery states

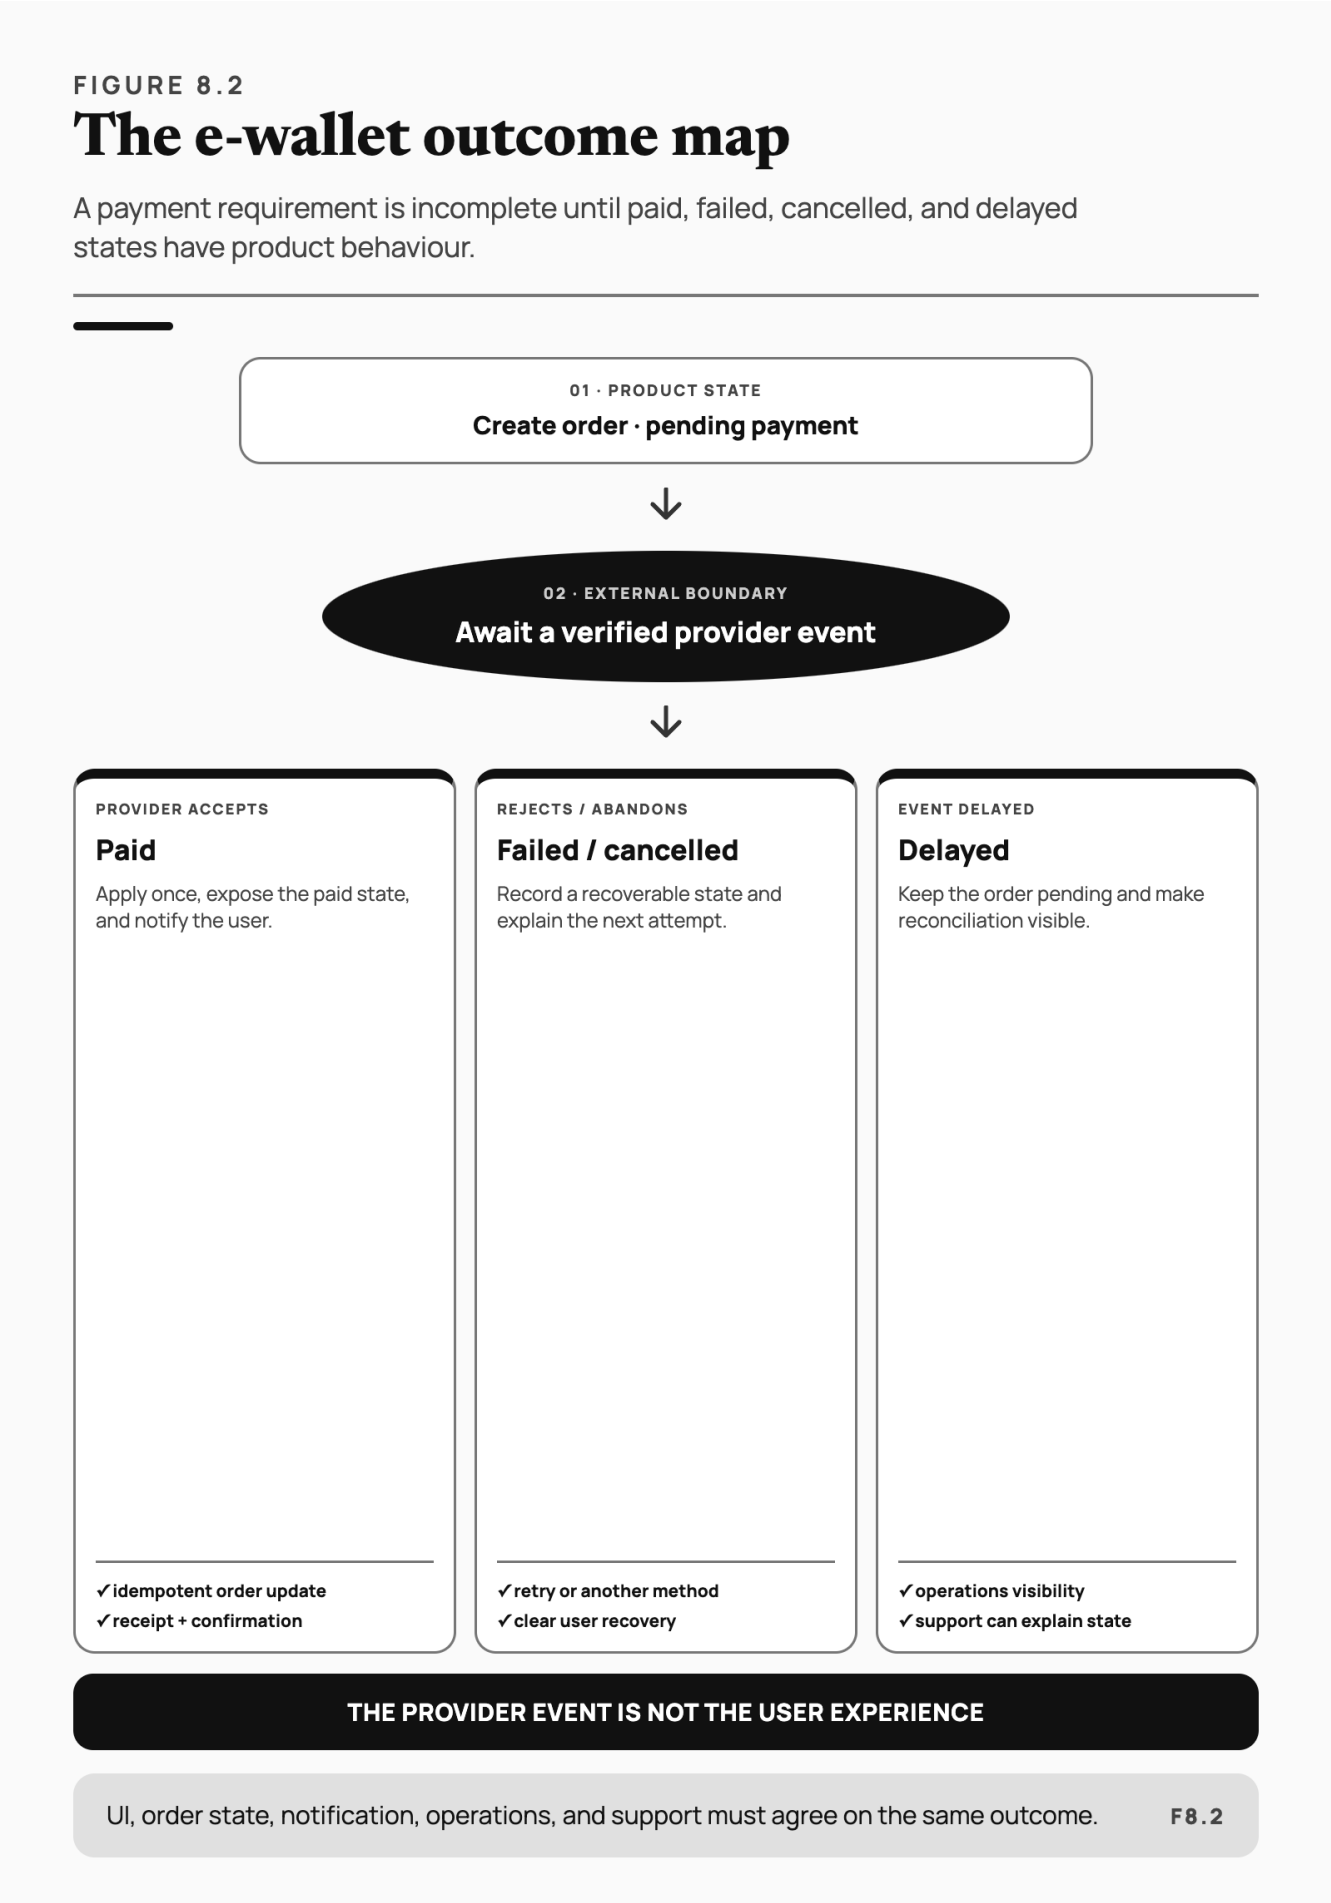

The e-wallet flow makes pending, failed, duplicate-safe, and support-visible behavior part of the requirement rather than late implementation detail.

Figure 9.1. Token-component-pattern-product hierarchy

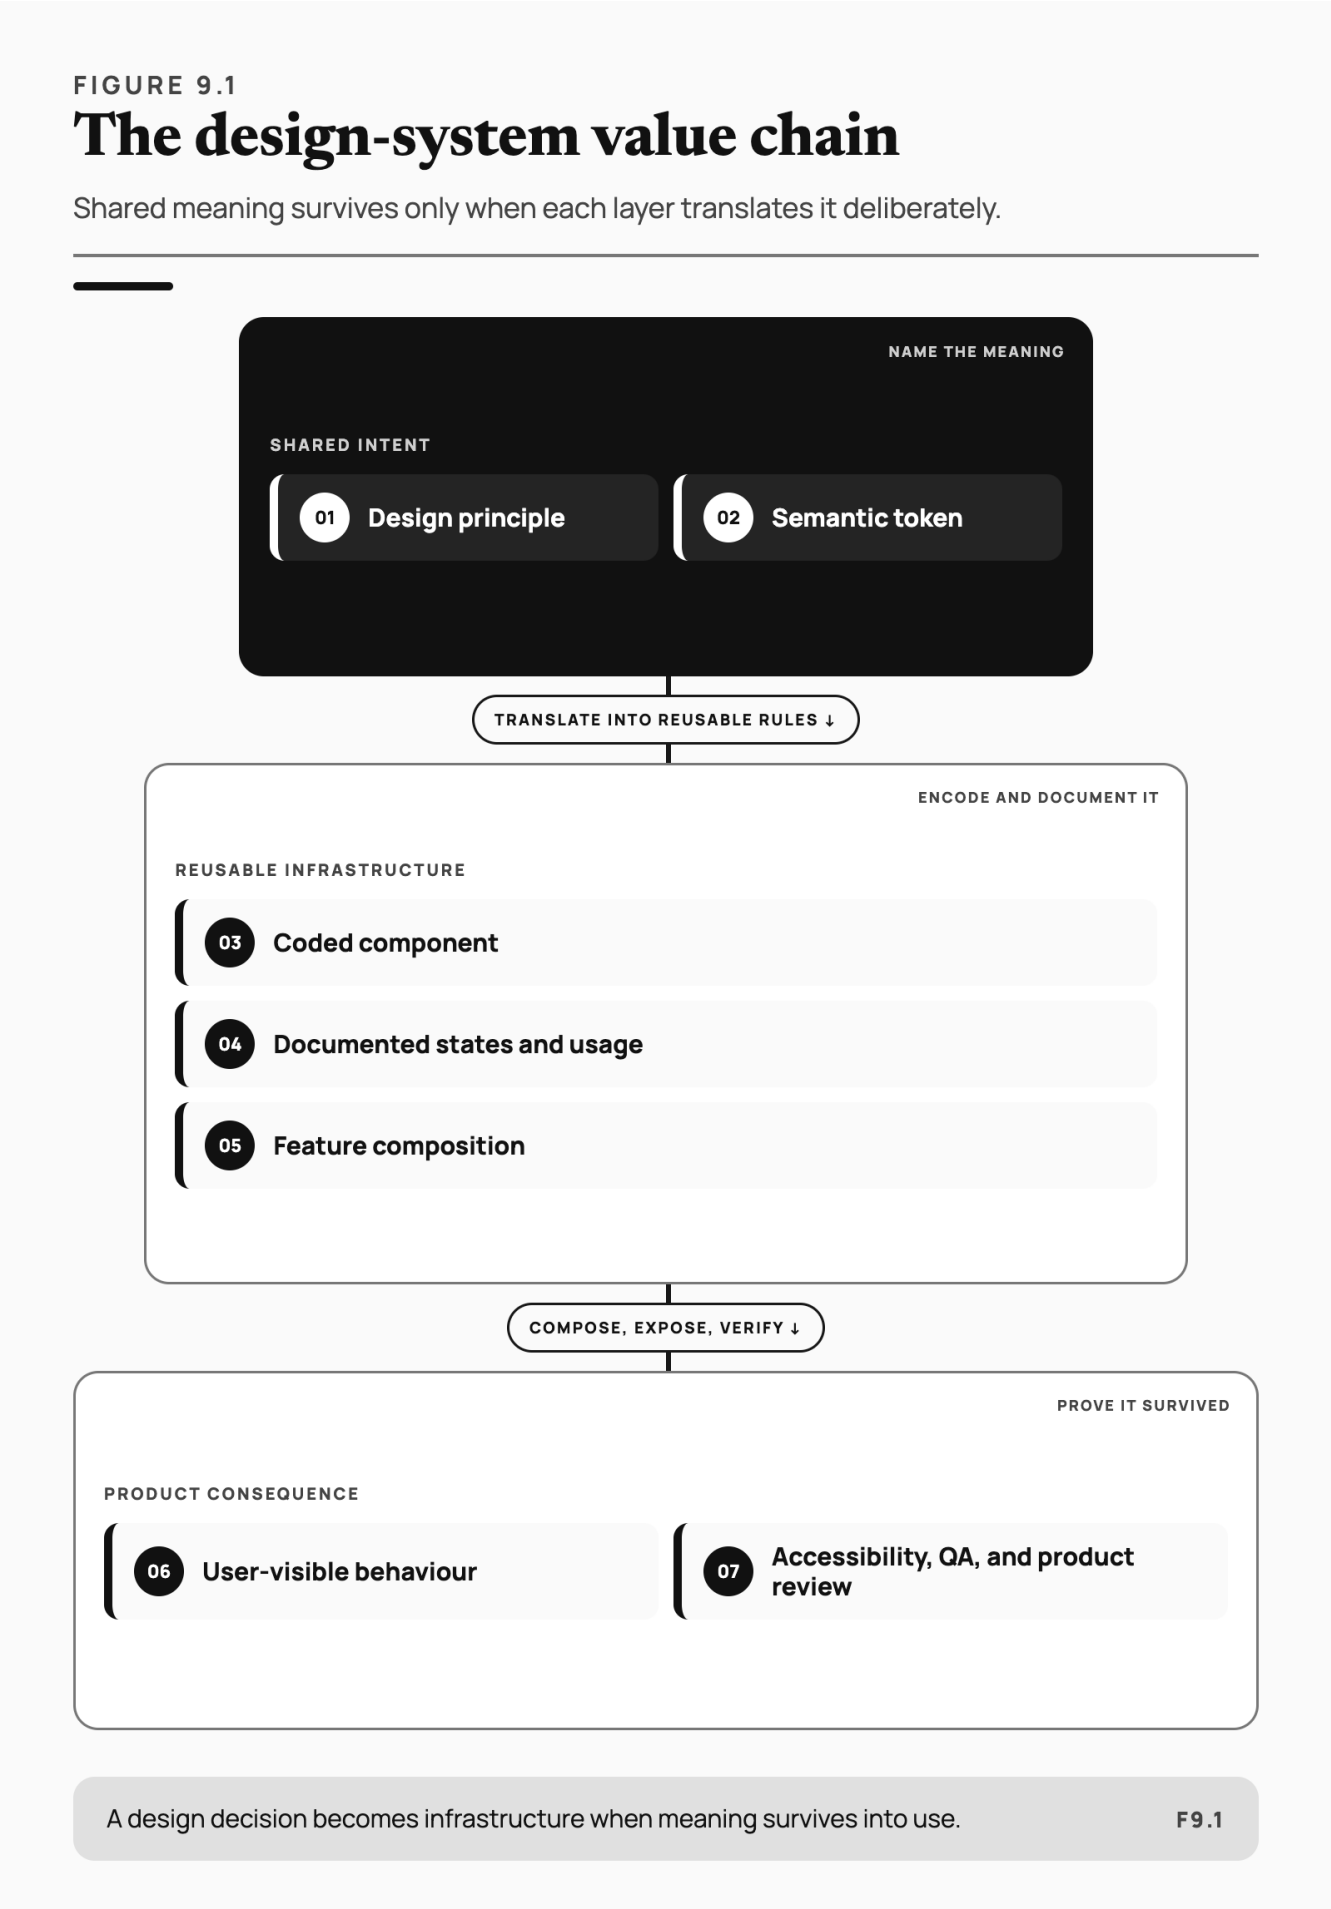

A design decision becomes infrastructure only when meaning survives from principle and token through code, states, composition, and review.

Figure 11.1. Onboarding context decision boundary

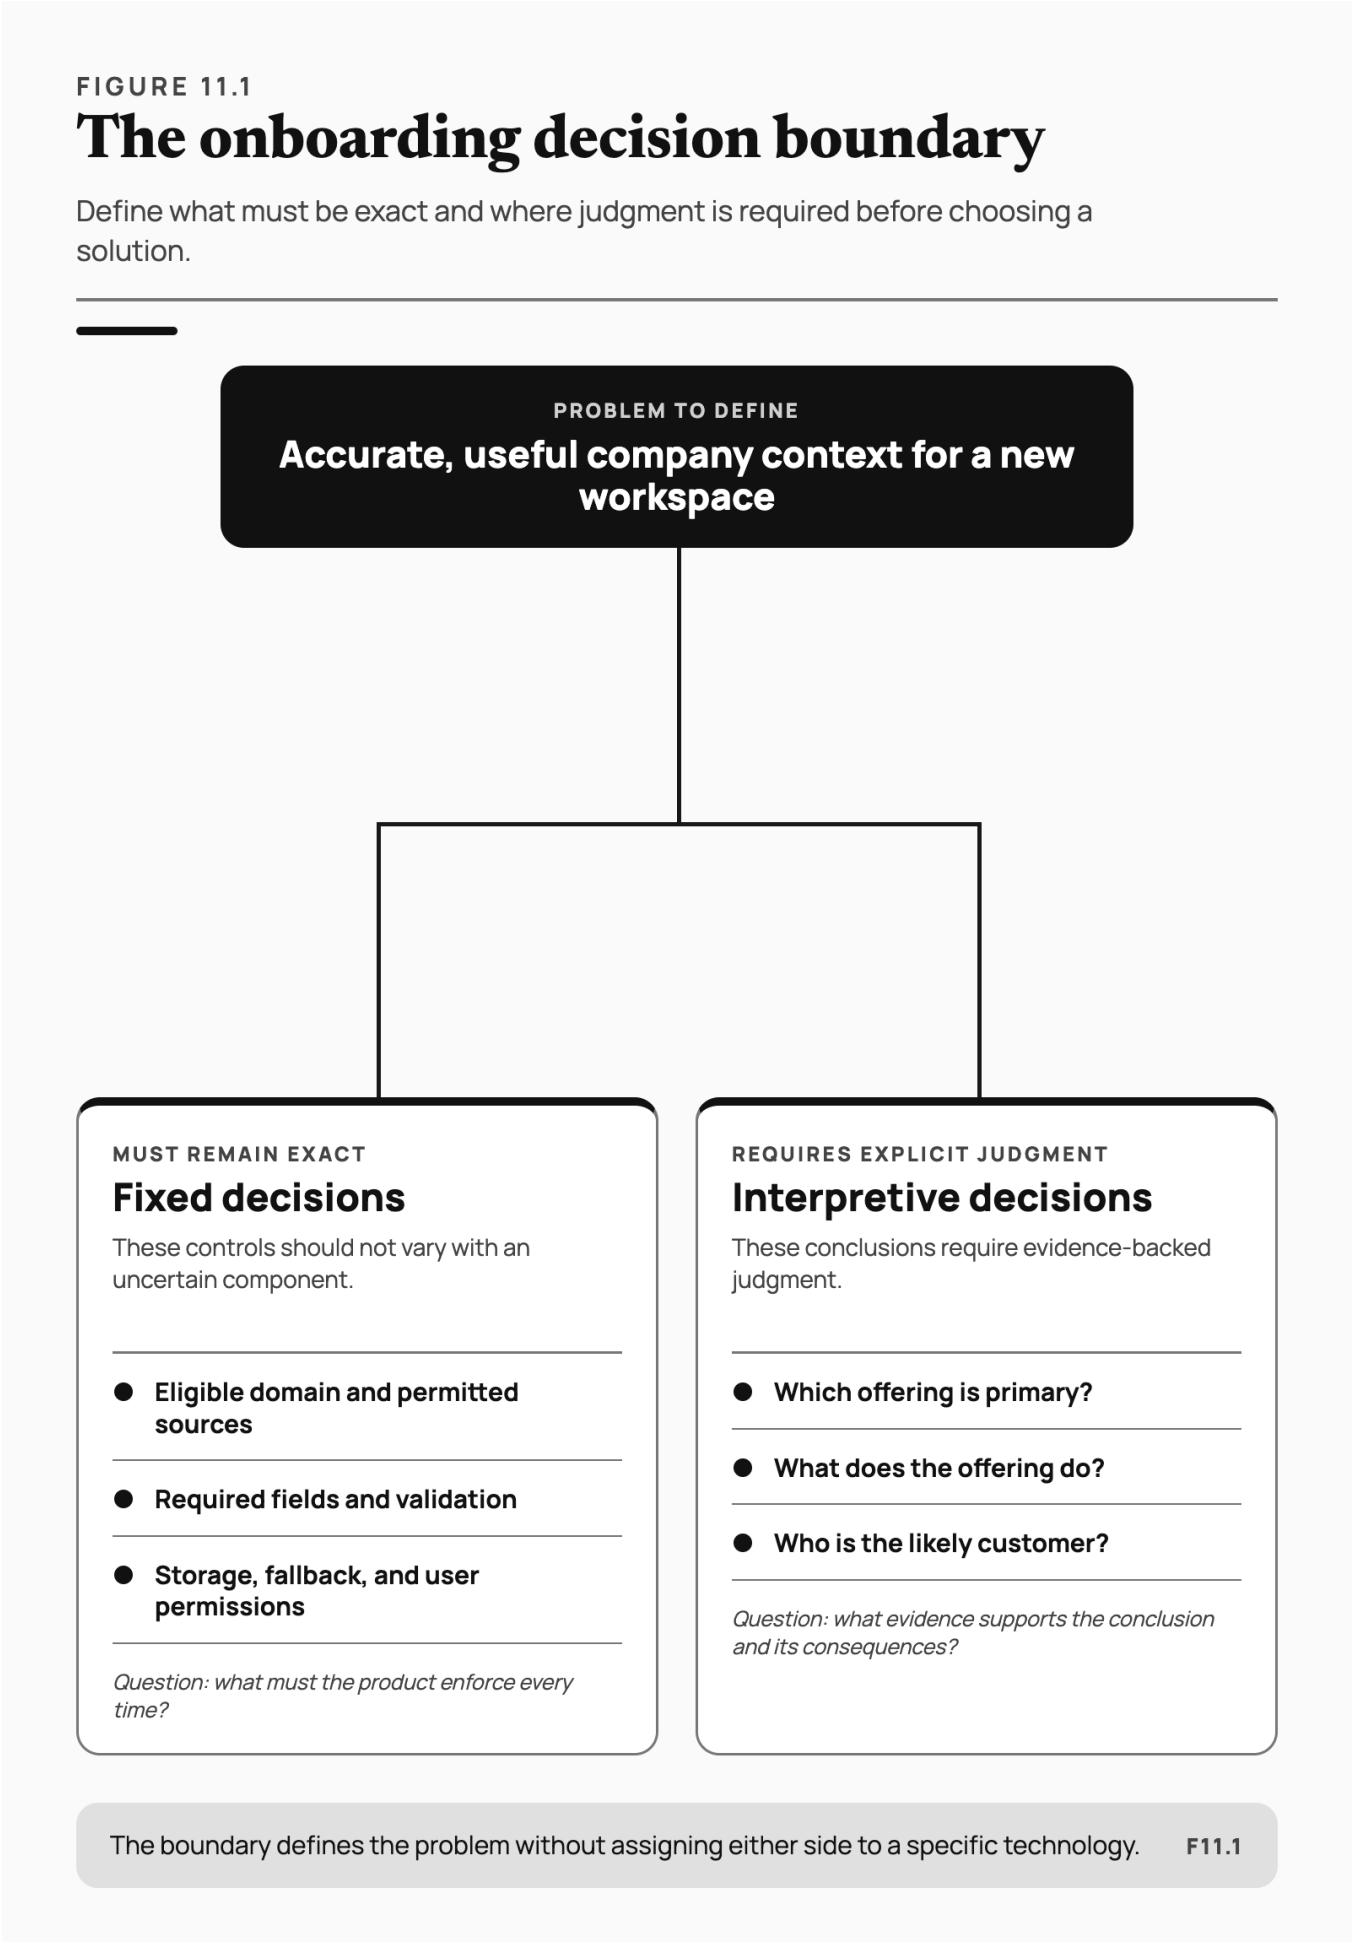

Separating fixed decisions from interpretive ones defines the problem boundary without choosing the implementation in advance.

Figure 12.1. Heatseeker hybrid solution shape

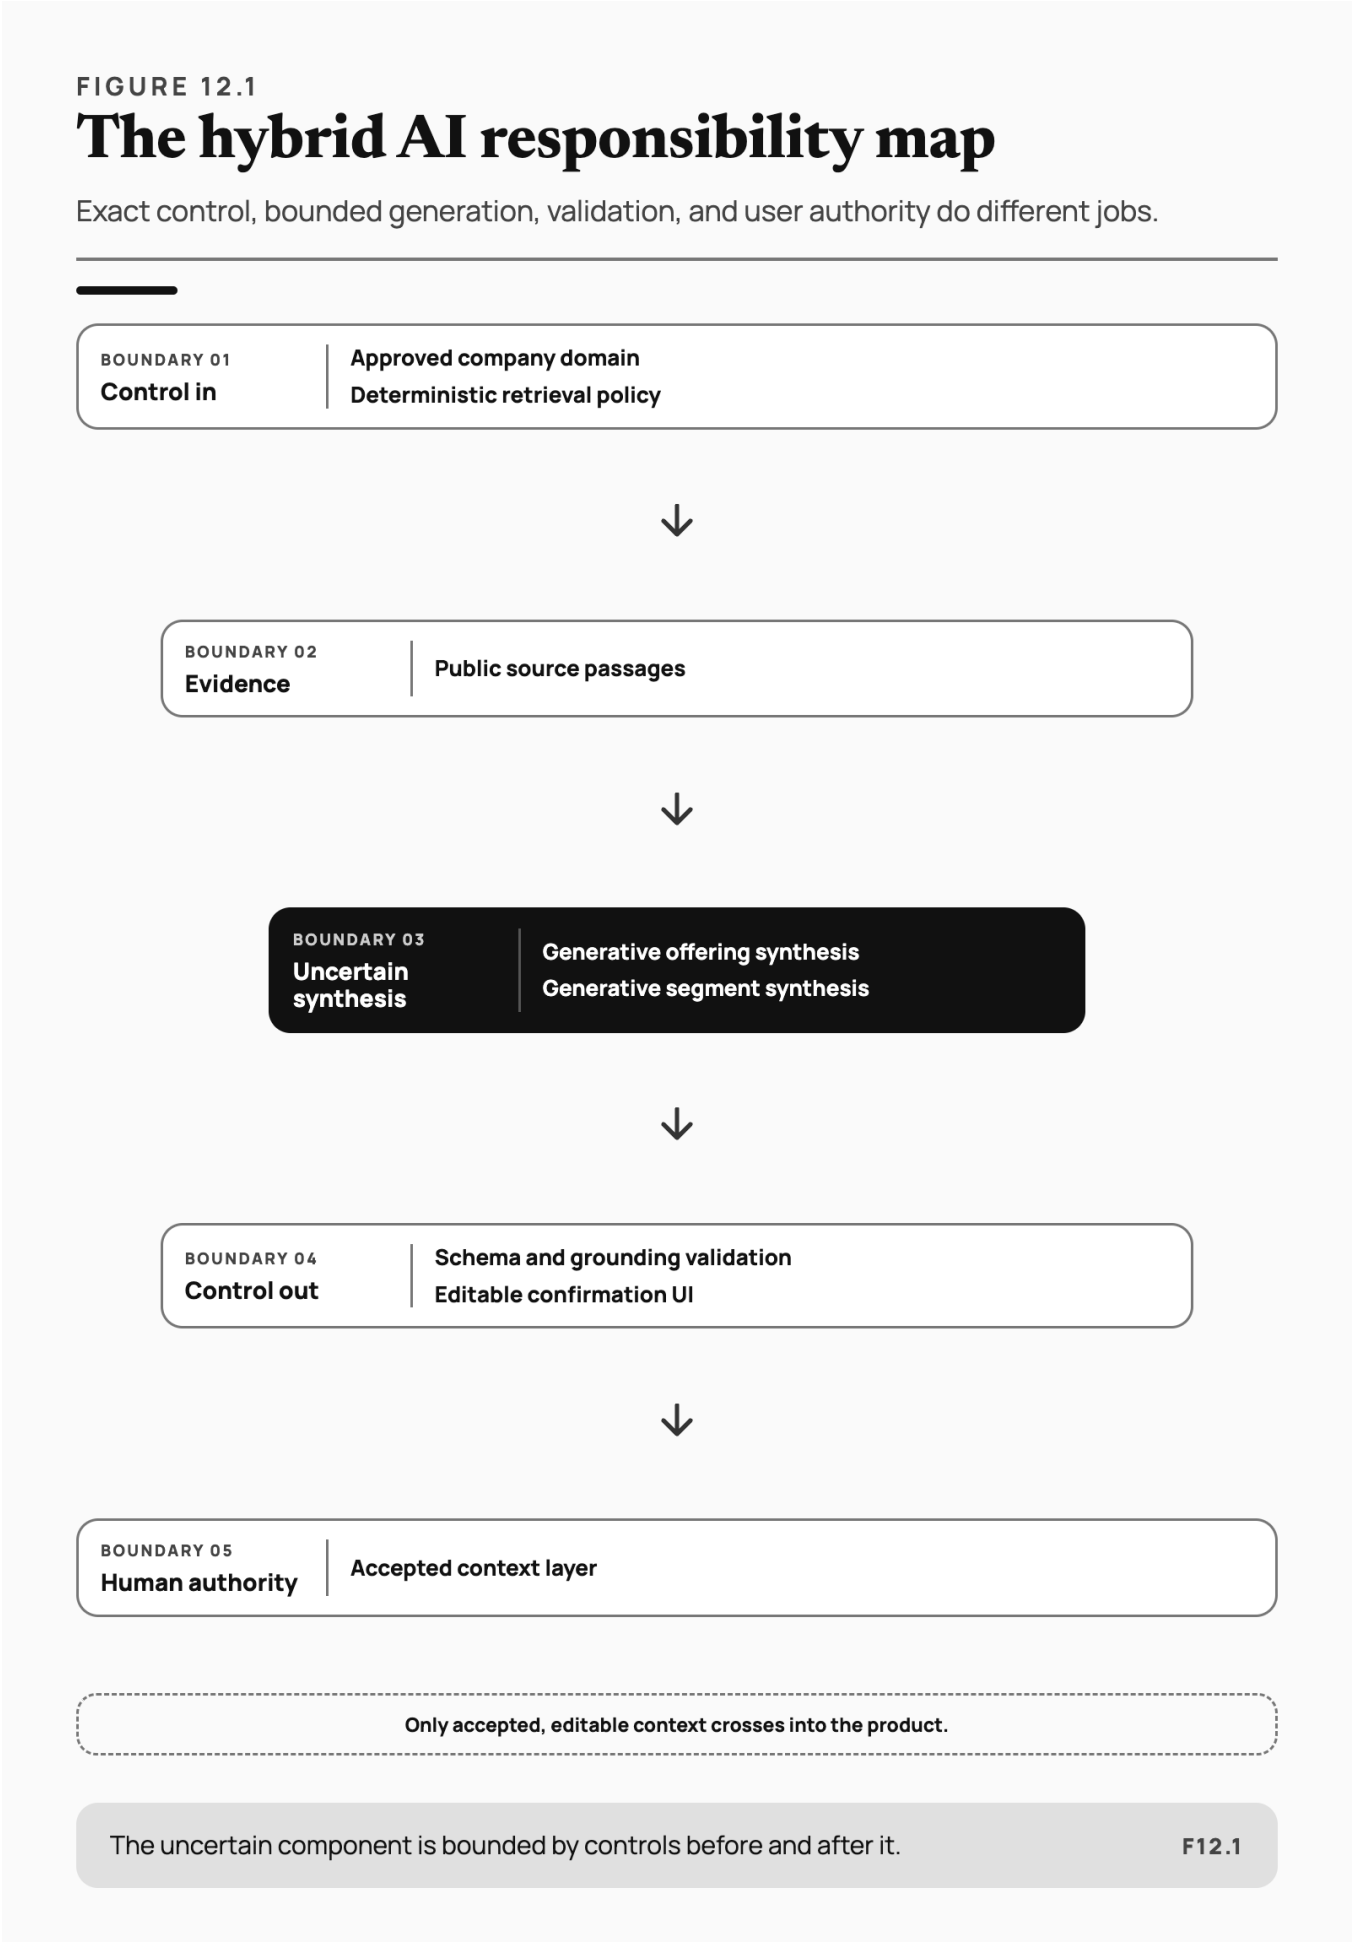

A hybrid design assigns exact control, bounded interpretation, validation, and user authority to different components.

Figure 12.2. Four production contracts around a model

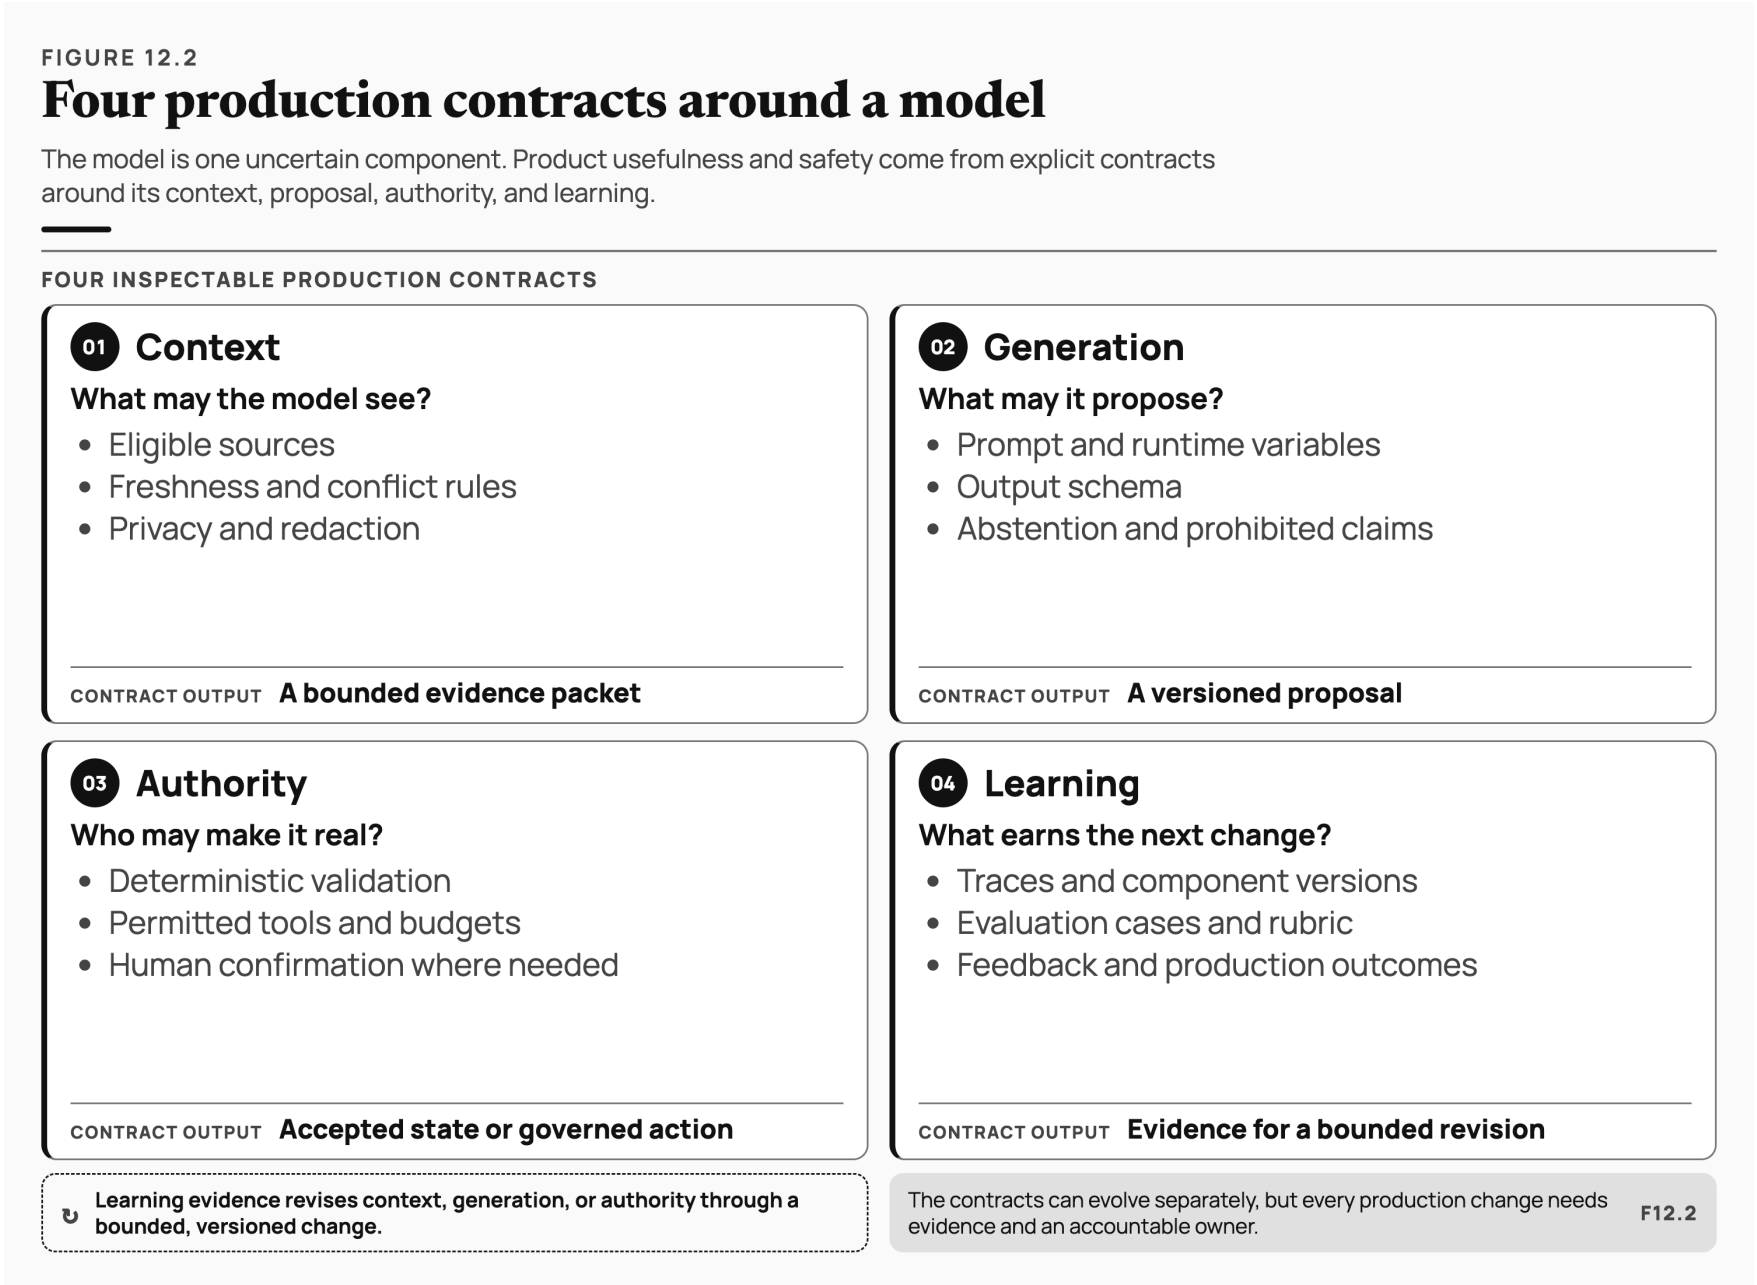

A production model feature is governed through four contracts: what context is eligible, what the model may return, who may authorise action, and how evidence changes the system.

Figure 13.1. Competitor-ad evaluation pipeline

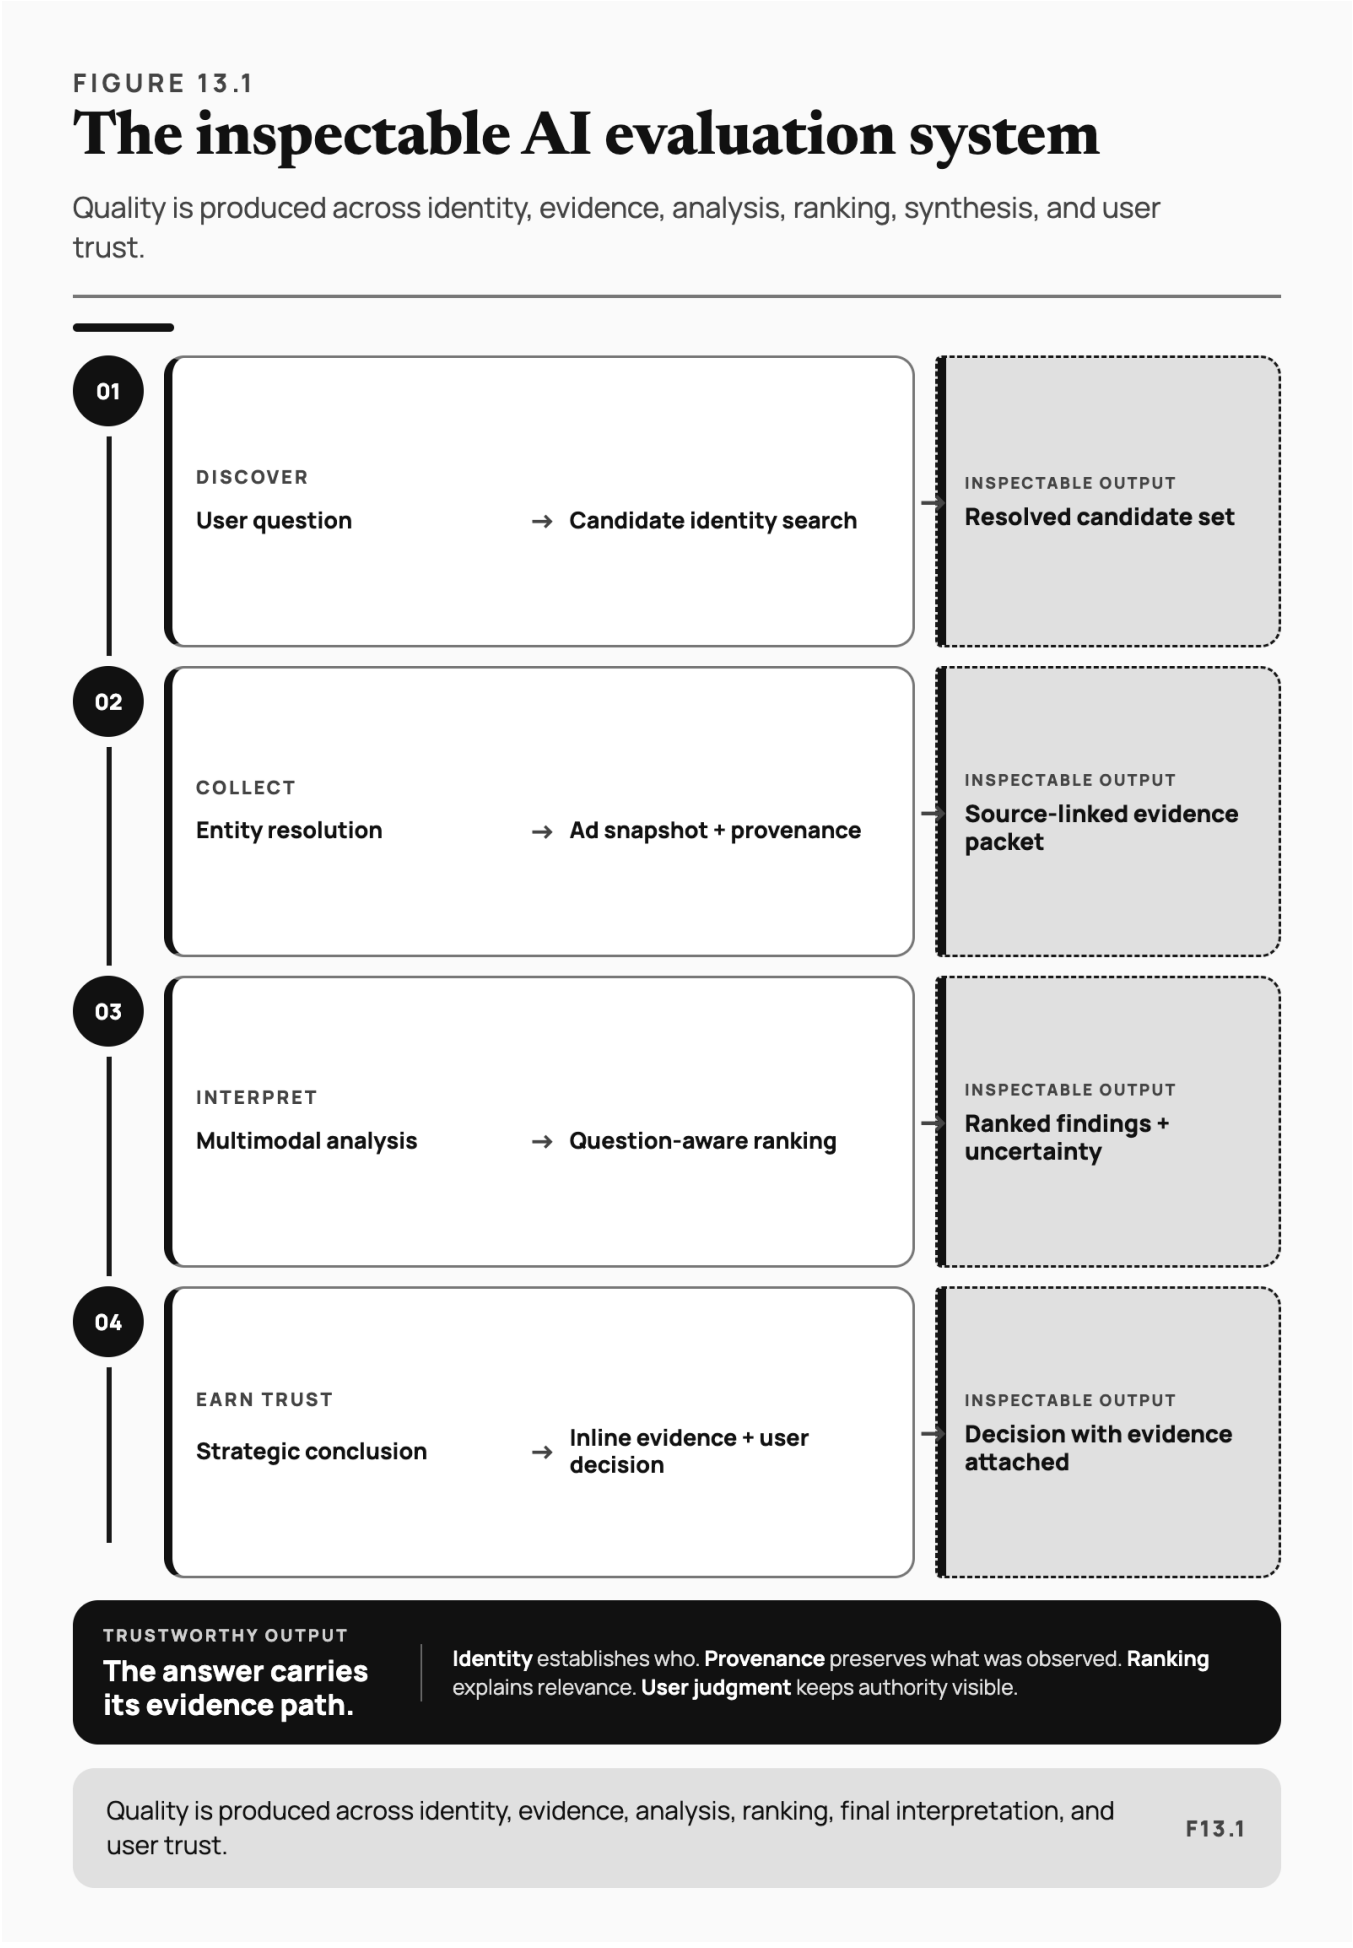

Final-answer quality depends on identity, evidence, analysis, ranking, final interpretation, and UI trust—not only on the generation model.

Figure 14.1. Human-directed Heatseeker calibration loop

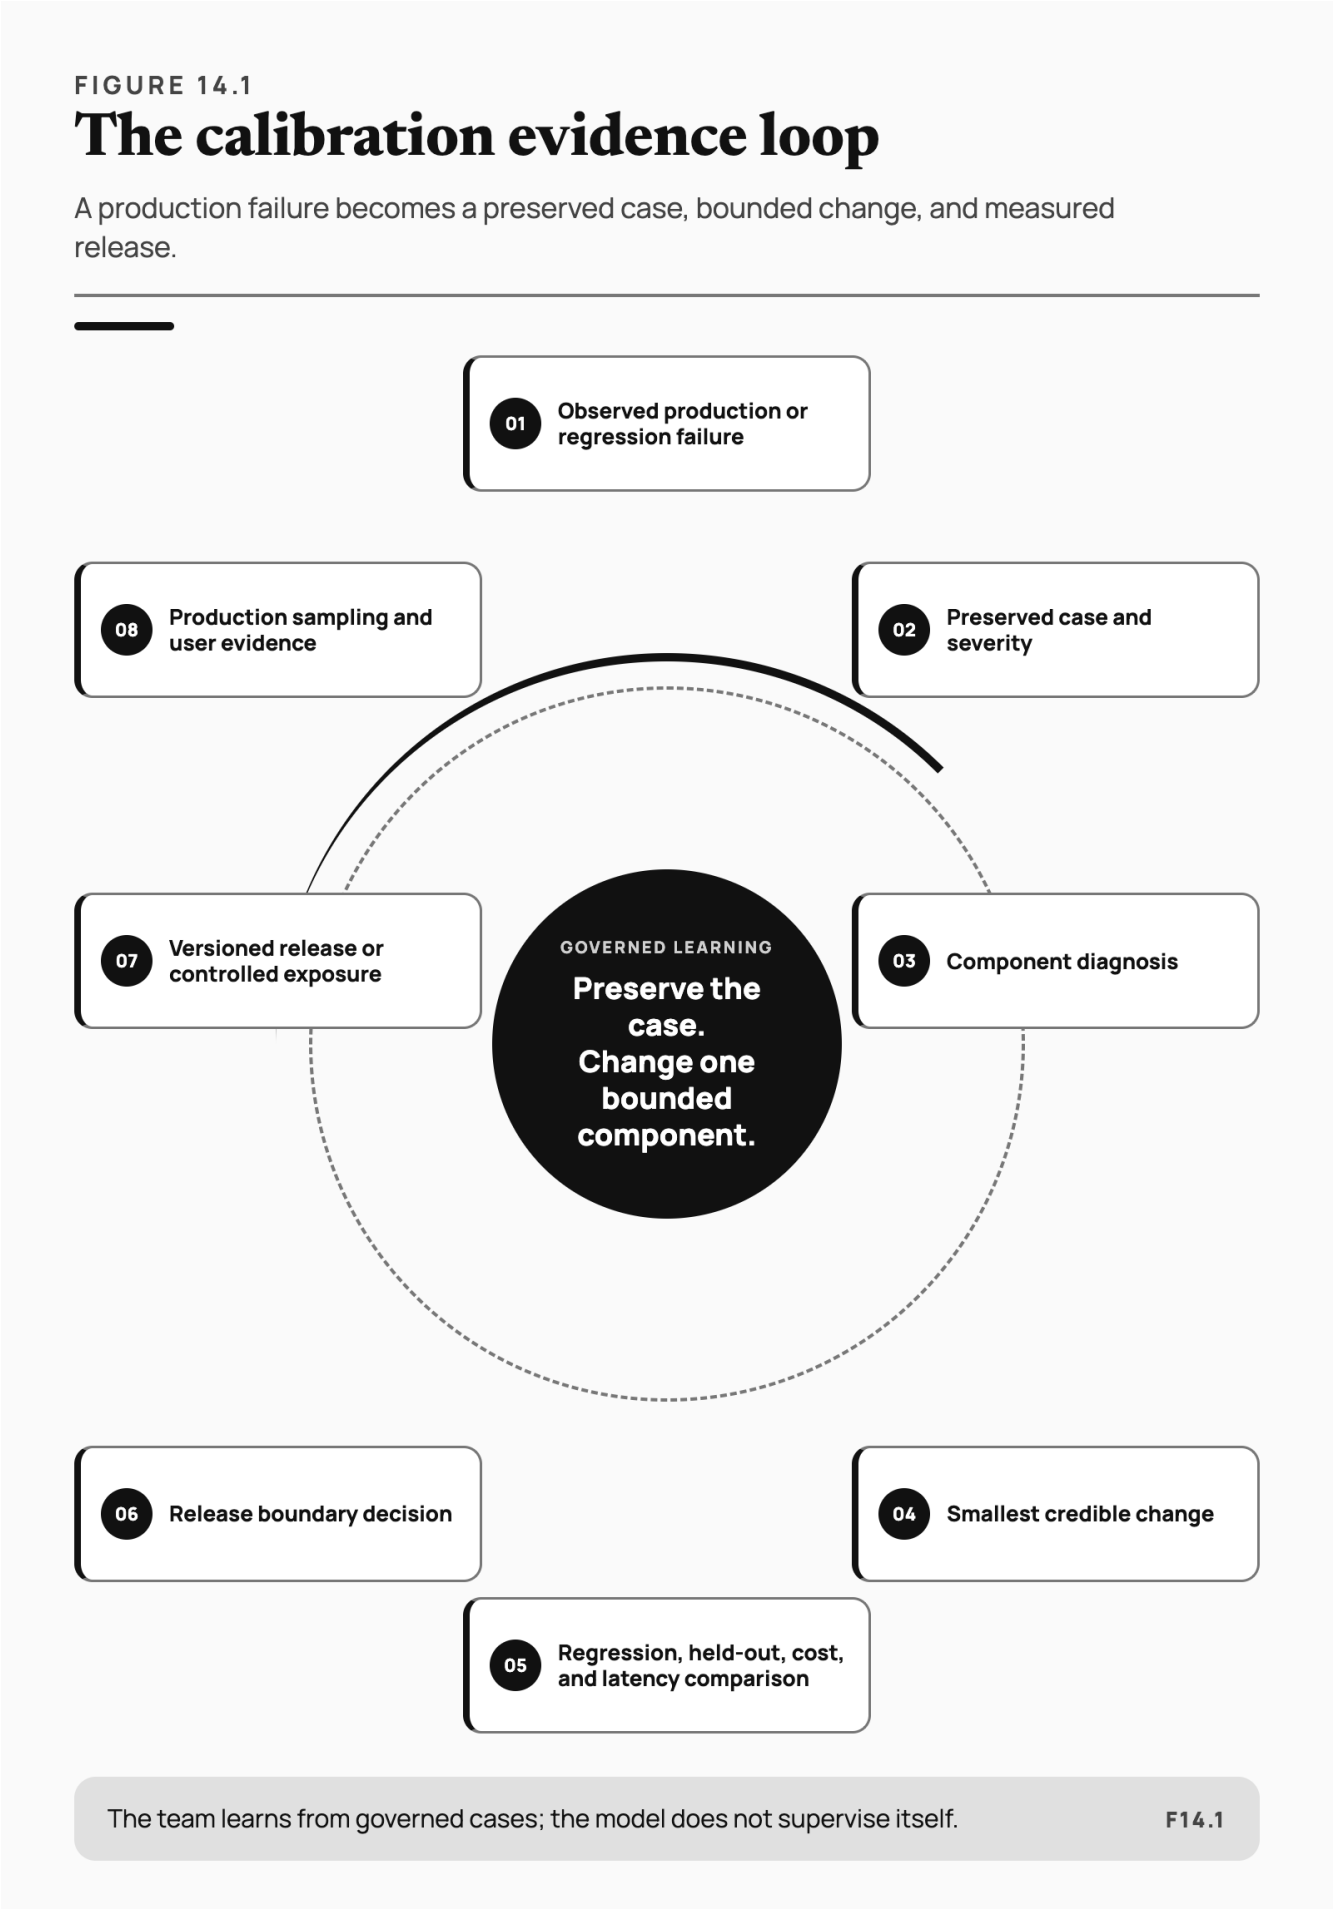

Calibration is a human-directed quality loop: an observed failure becomes a preserved case, bounded change, regression-tested release, and new production evidence.

Figure 15.1. Bounded competitor-analysis loop

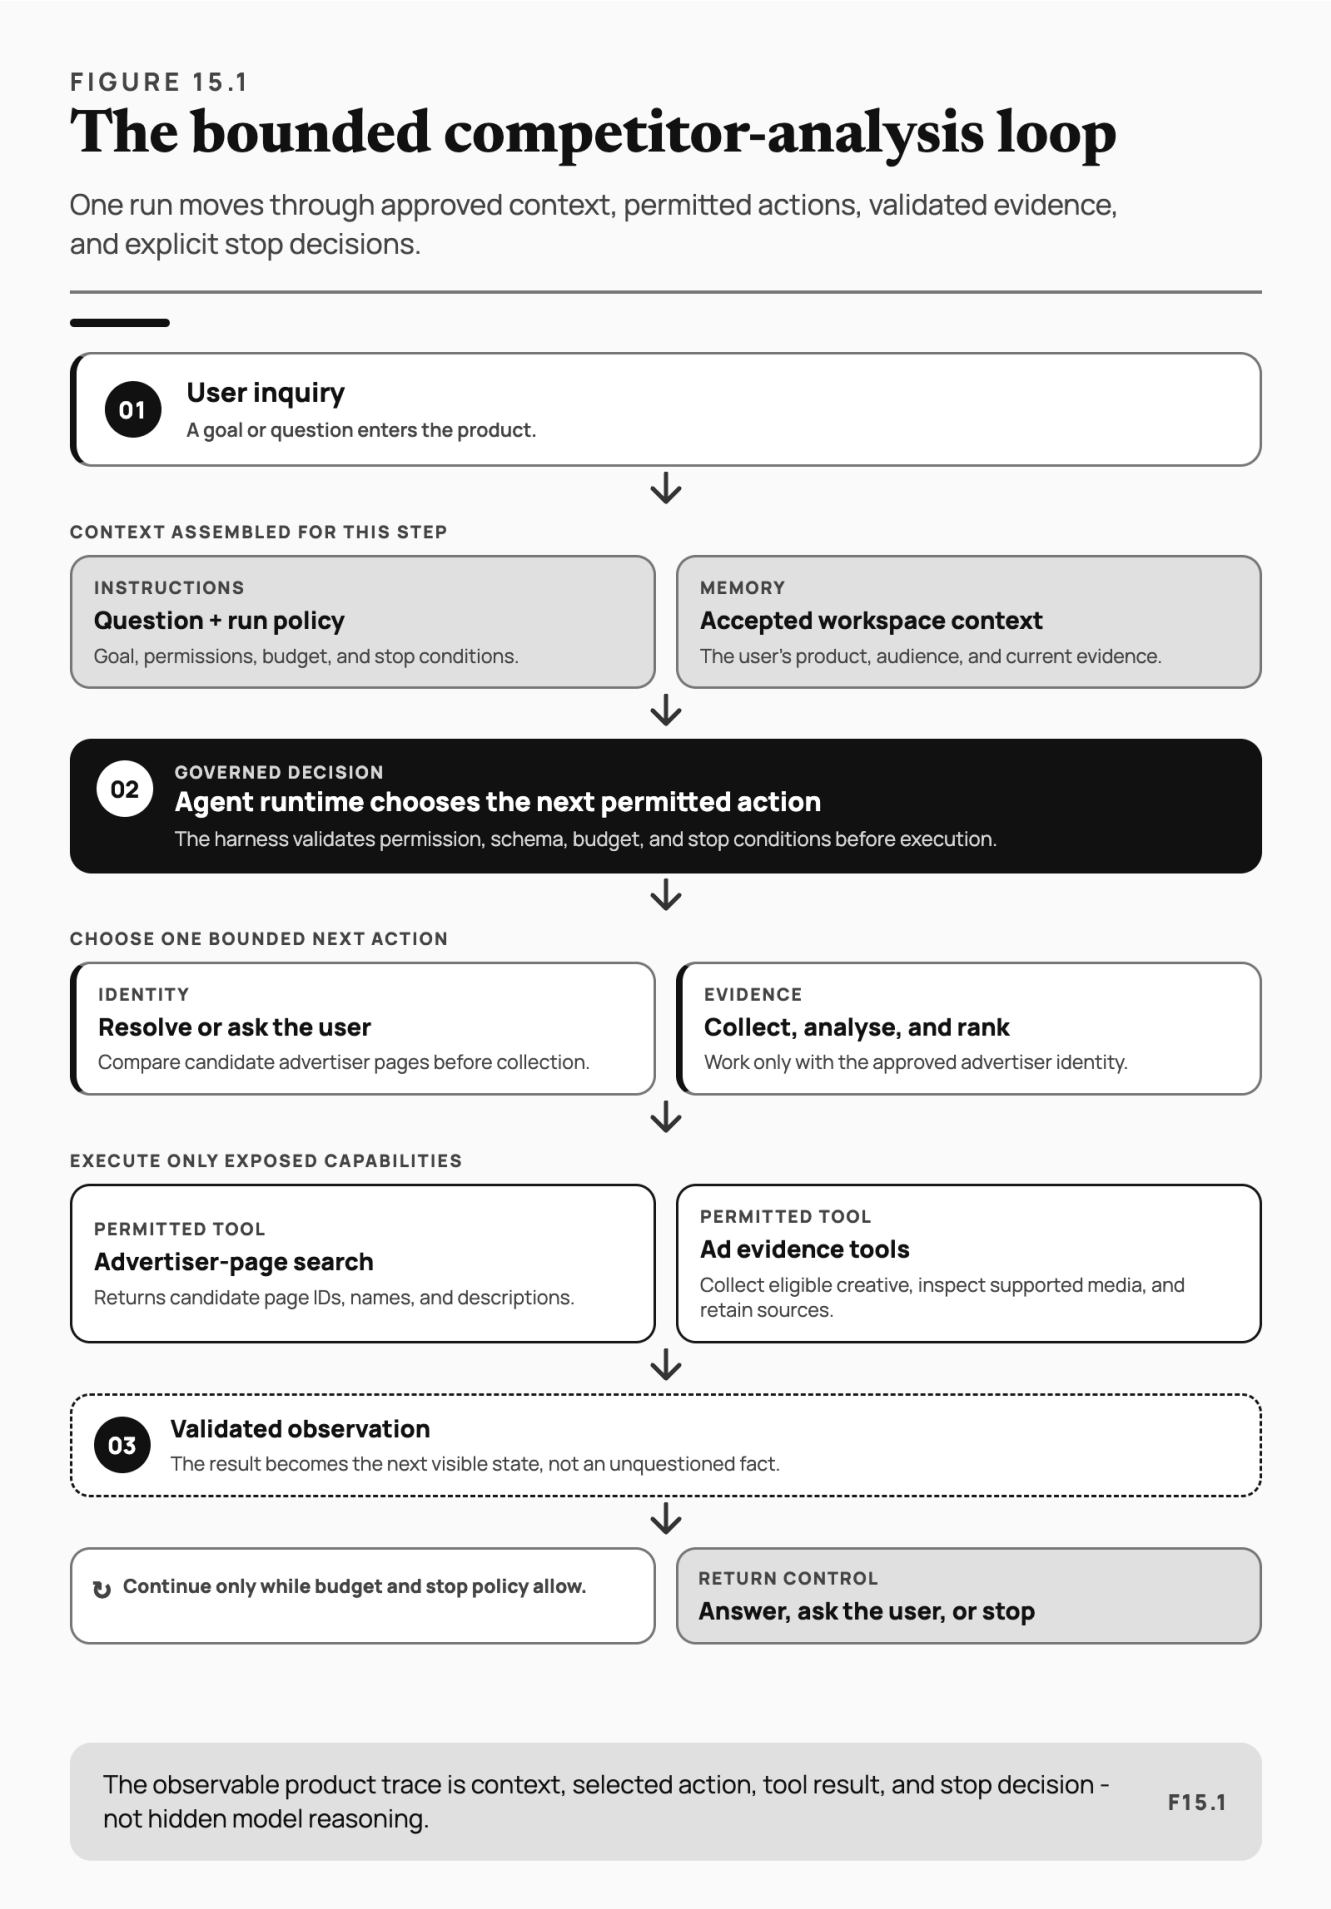

The observable competitor-analysis loop: the application assembles governed context, permits a bounded next action, validates the result, and either continues within its limits or returns control to the user.

Figure 16.1. Agent operating contract

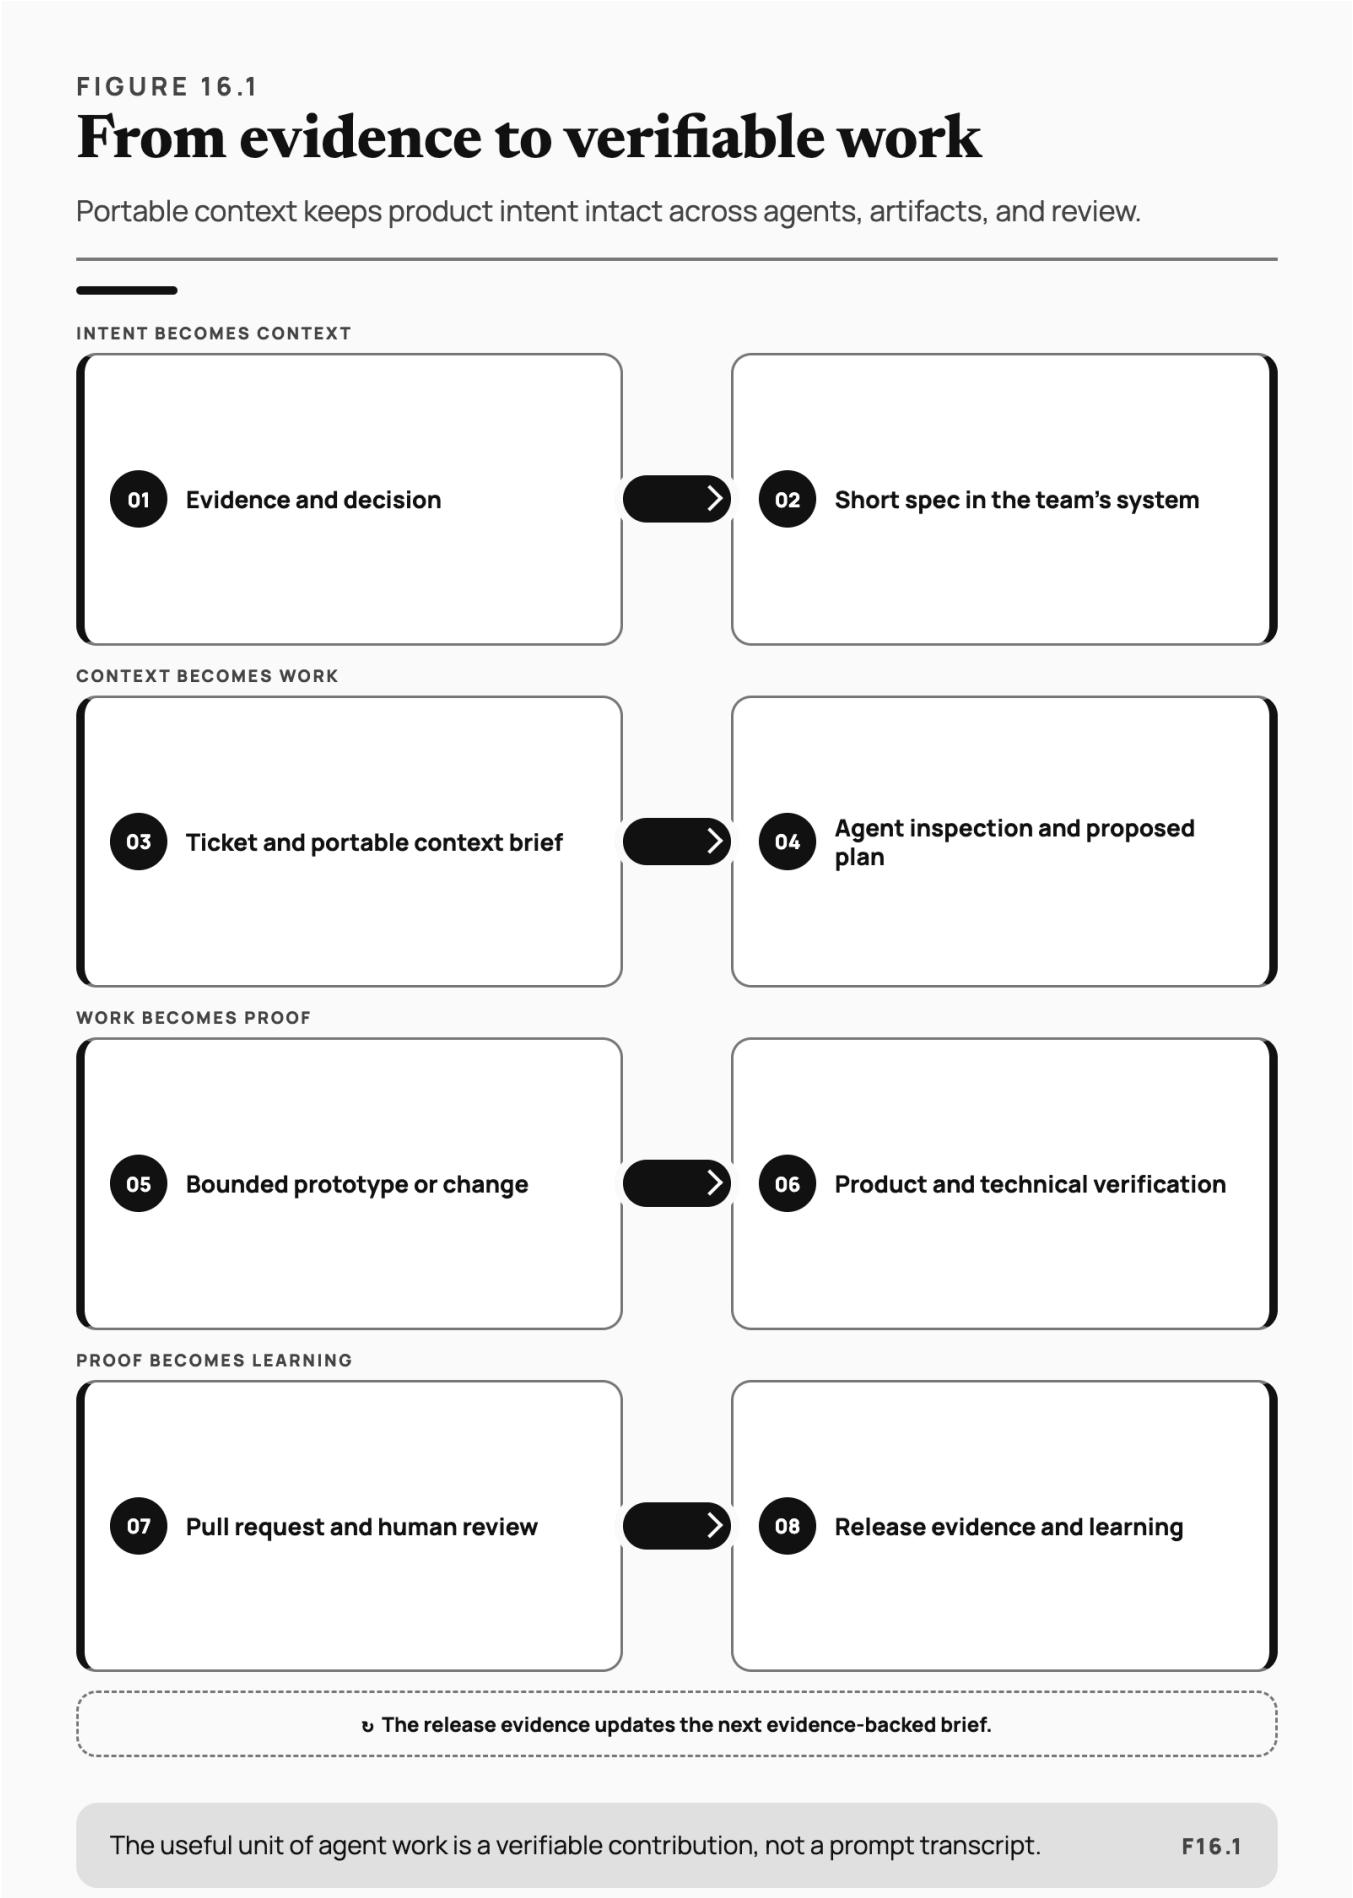

Agent context should preserve the evidence-to-decision thread and end in human-verifiable work, not an isolated prompt transcript.

Figure 17.1. Repo as product surface

Repo orientation connects the user promise to implementation, checks, environments, and accountable reviewers; product context can explain rationale, but it does not replace runtime verification or human ownership.

Figure 17.2. Noted cover-image upload

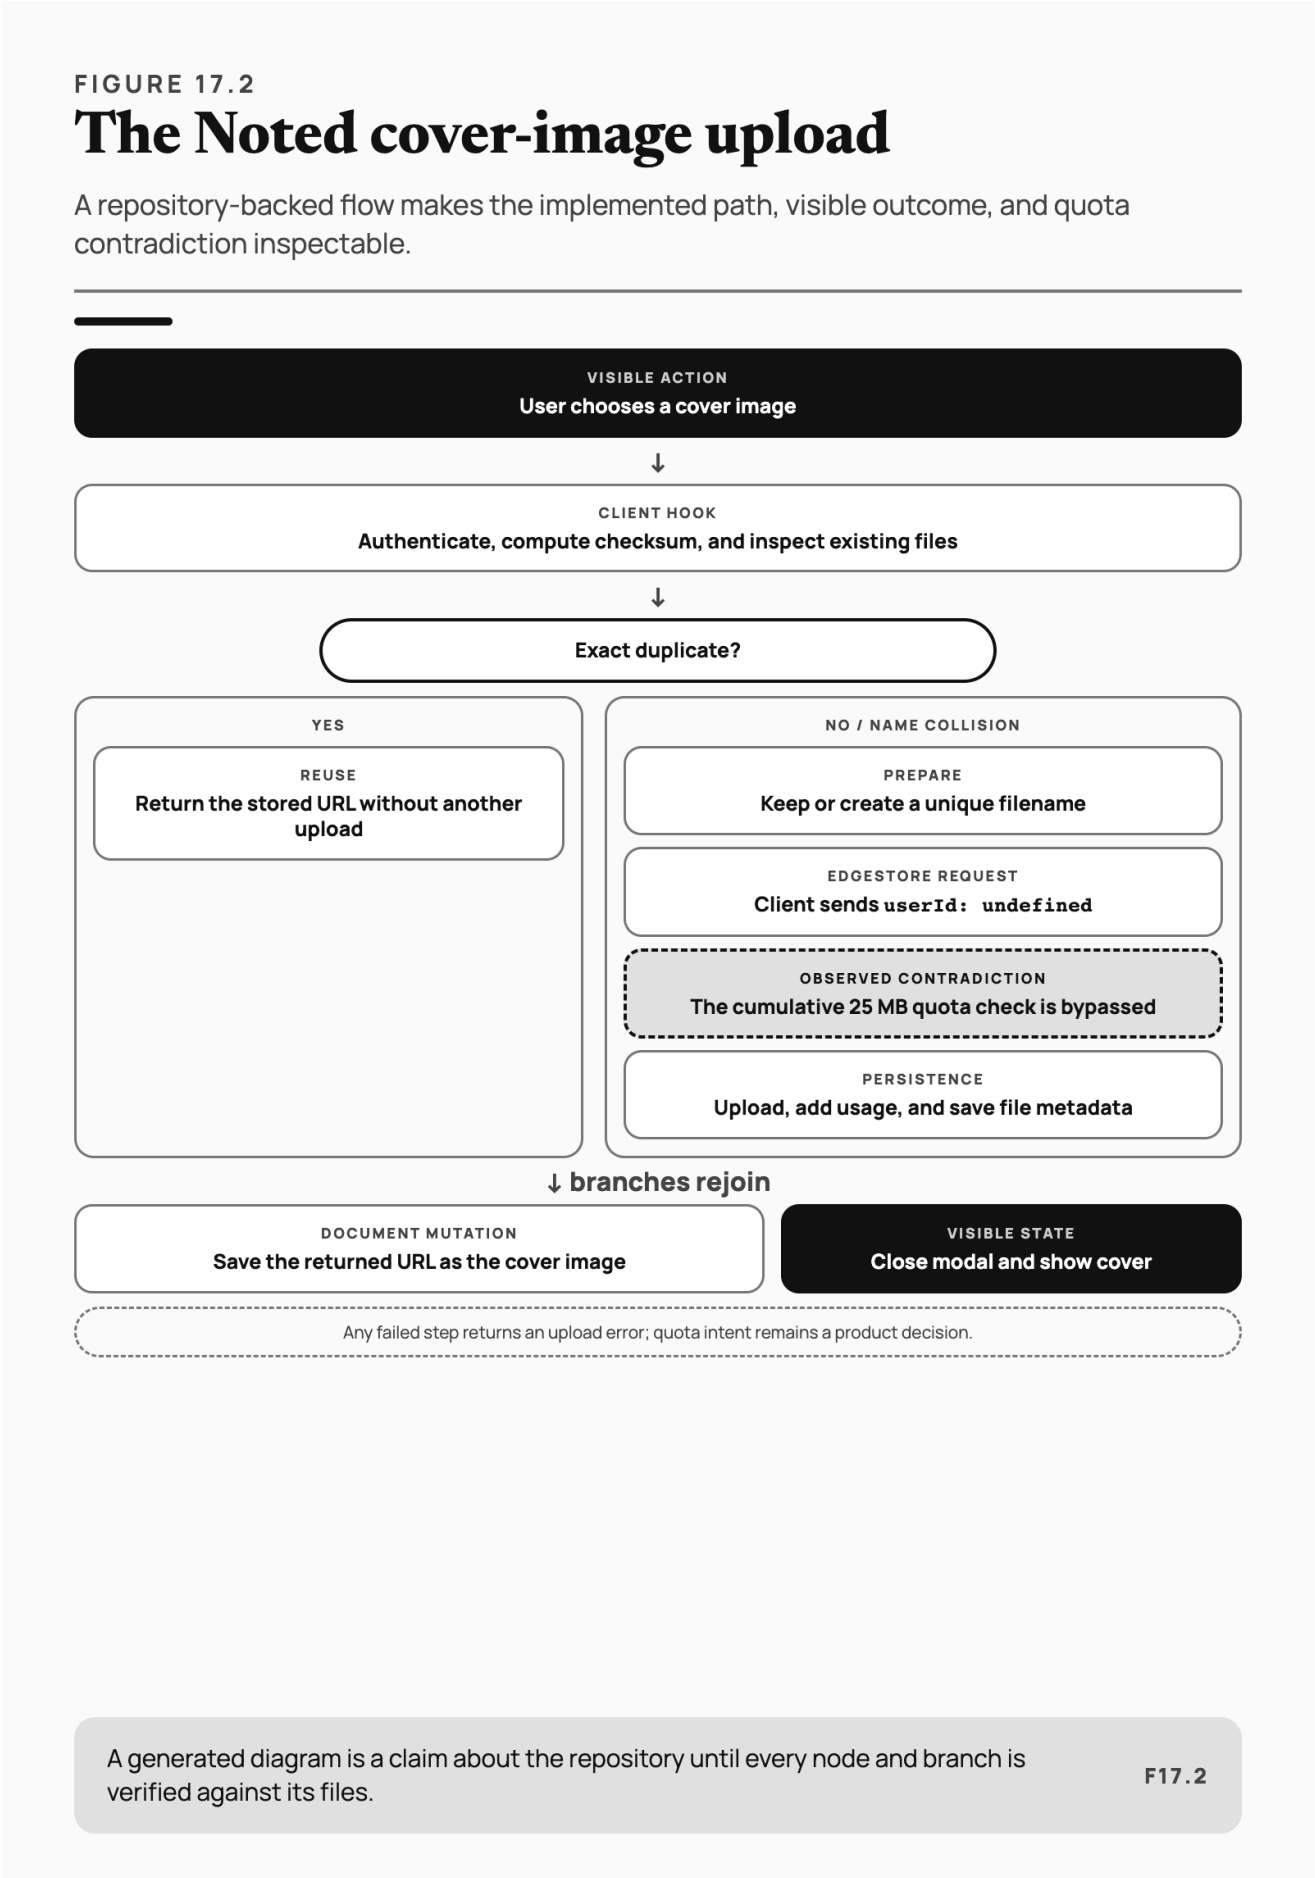

A repository-backed Mermaid flow for Noted's cover-image upload at commit `985ad95`. The branch makes the central contradiction visible: the route has a cumulative-quota check, but the current client path sends no user ID and therefore bypasses it. Every node should be verified against the cited files before the diagram is treated as product truth.

Figure 18.1. Evidence operations continuity

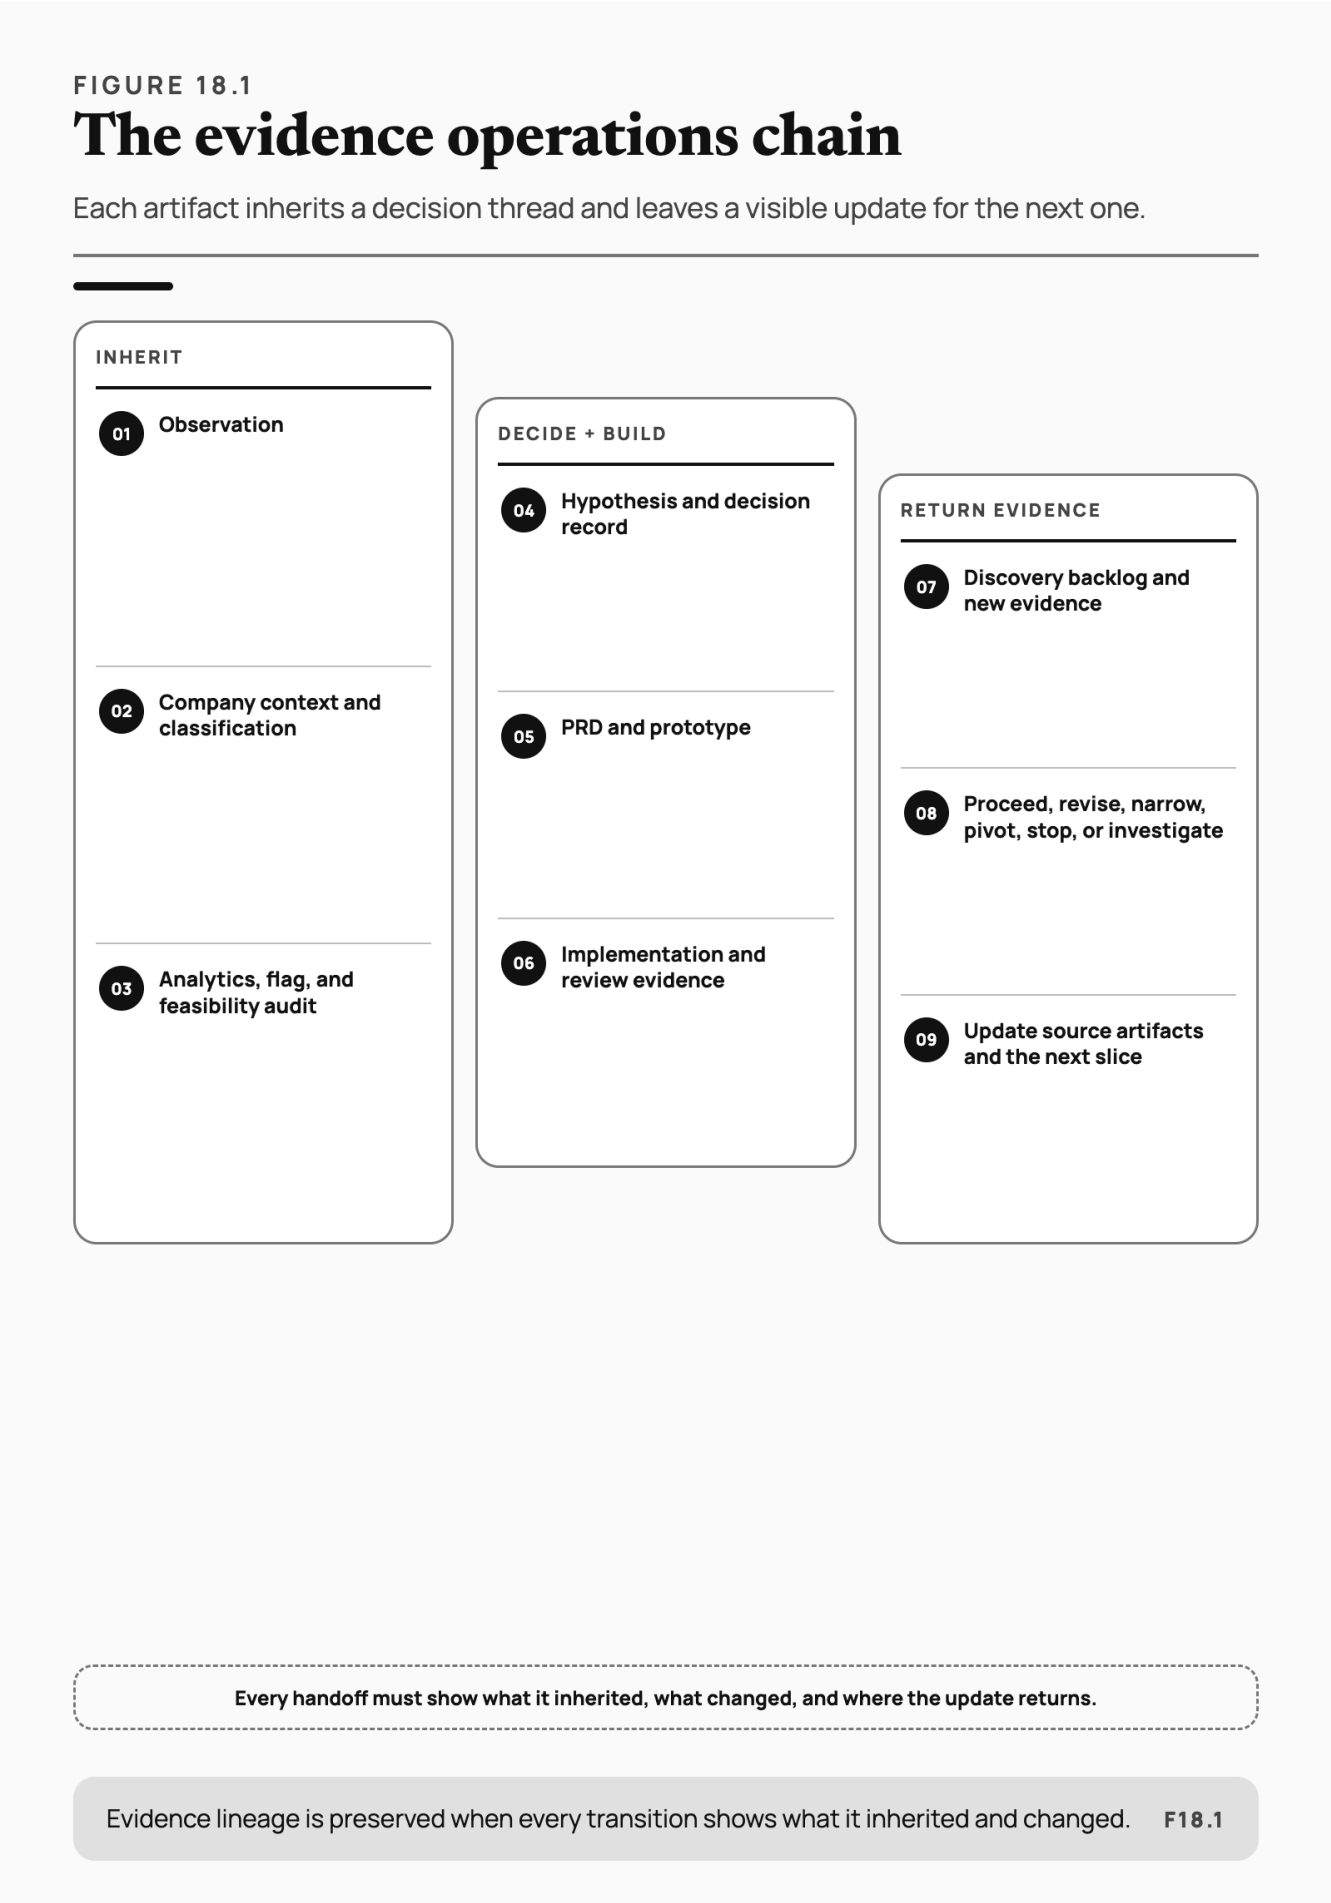

Evidence operations preserve lineage by carrying each observation through classification, product artifacts, implementation evidence, and an explicit decision that updates the next source of truth.

Figure 19.1. Prototype fidelity ladder

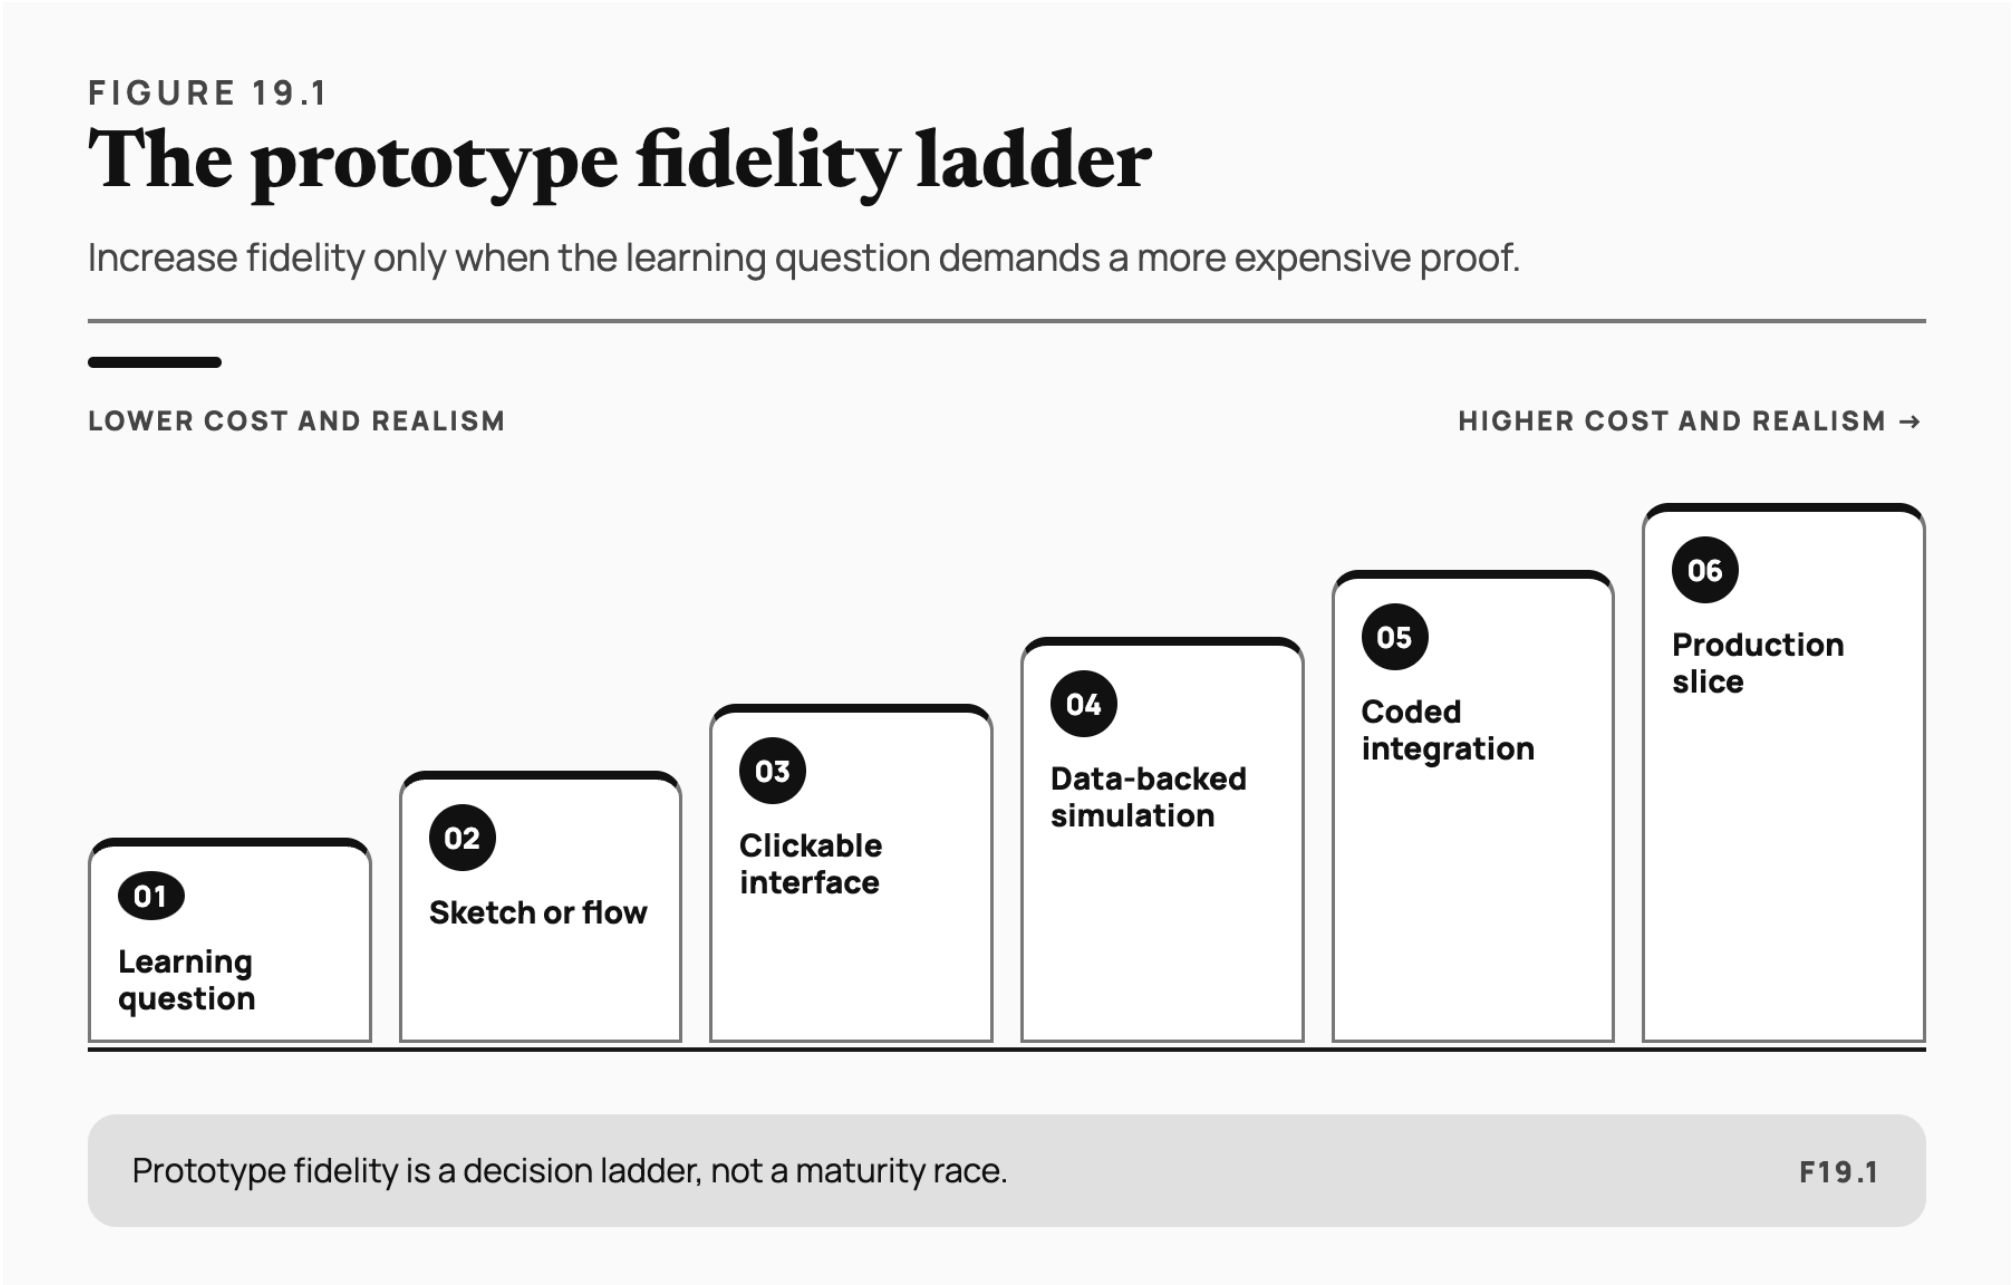

Prototype fidelity is a decision ladder, not a maturity race: move only as far as the learning question requires, and treat production exposure as a separate commitment.

Figure 20.1. Building is not shipping

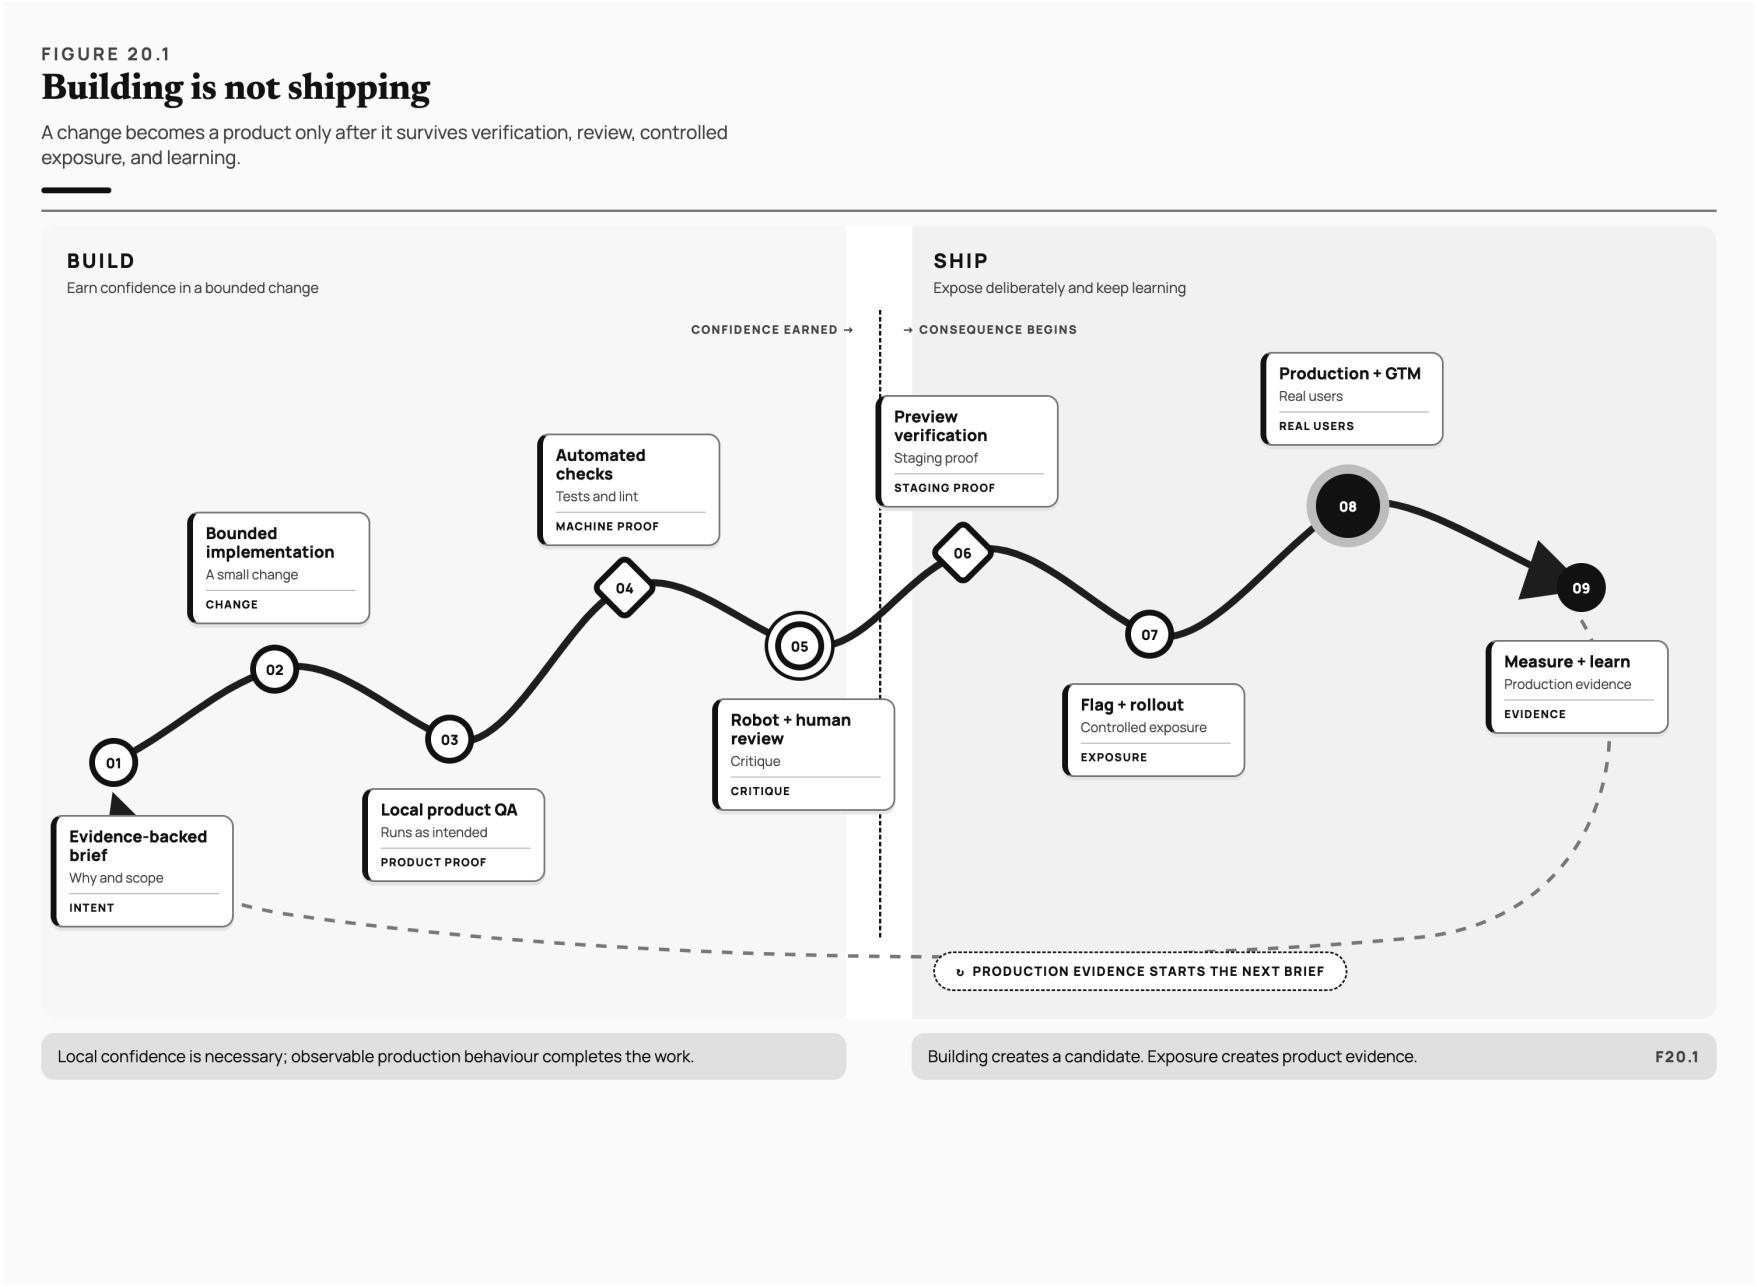

Building becomes shipping only after the change survives verification, review, controlled exposure, and measurement.